Commitment Floors for Tipping-Point Commons: Escaping Nash Traps in Multi-Agent RL with Safety-Constrained Moral Architectures

Anonymous Author(s)

Abstract

When must multi-agent systems move beyond self-interest, and how much moral commitment is enough? We formalize this question through the class of **Tipping-Point Social Dilemmas** (TPSDs)—environments where non-linear resource dynamics create irreversible collapse thresholds—and identify a systematic failure mode we term the **Nash Trap**: individually rational policies converge to locally stable equilibria that guarantee collective extinction. We make three contributions: (C1) **Nash Trap Analysis**: We prove that the Nash Trap arises from a structural mismatch between individual payoff gradients and collective survival requirements, and demonstrate that this failure is *consistent across seven MARL paradigms*—REINFORCE (3 architectures), PPO (IPPO/MAPPO), IQL, QMIX, and LOLA—with a unified fair comparison protocol ($400+$ HP configurations, 100% trap rate). (C2) **Commitment Floor as Safety Specification**: We derive that the minimum commitment floor ϕ_1^* for survival is determined by the environment's recovery function near its tipping point, establishing $\phi_1^* = 1.0$ as a *stability-derived design specification* for severe TPSDs. A constrained policy optimization formulation independently converges to the same floor, confirming it is not an arbitrary heuristic. (C3) **Generalization and Robustness**: We validate the Nash Trap and commitment floor across multiple environment classes (common-pool resources, spatial grid dilemmas, Cleanup-style commons), 36 shock parameter regimes, adversarial conditions (30% Byzantine agents), and at scale (N=100). Commitment floors simultaneously improve fairness (Gini: $0.064 \rightarrow 0.000$, worst-case return: $+87\%$). Together, these results establish a **Moral Commitment Spectrum**: the degree of architectural commitment required for collective survival is a monotonic function of environmental severity.

1 Introduction

Rational Choice Theory defines human behavior as utility maximization under constraints. Sen [1977] critiqued this as the "Rational Fool," arguing that genuine moral behavior requires *commitment*—acting against immediate self-interest—mediated through *meta-rankings*: preferences over preference orderings. This philosophical insight gains practical relevance as multi-agent reinforcement learning (MARL) systems are increasingly applied to domains that may exhibit **non-linear dynamics and tipping points** [Lenton et al., 2008], where failure to cooperate can trigger irreversible collapse.

Can self-interested agents discover cooperative strategies through reward maximization alone, or must cooperation be architecturally encoded? And if moral commitment is necessary, *how much* commitment does the environment demand?

We address these questions through a systematic investigation spanning environments of increasing severity, yielding three contributions:

(C1) **Nash Trap Analysis**: We define the class of Tipping-Point Social Dilemmas (Definition 2), prove a drift condition on resource dynamics (Lemma 2), derive finite-horizon survival bounds (Lemma 3), and establish the Nash Trap as a locally stable fixed point (Proposition 4) where individually rational cooperation is insufficient for collective survival. We demonstrate that this failure is *consistent across seven MARL paradigms*: REINFORCE (3 architectures), PPO (IPPO/MAPPO), IQL, QMIX [Rashid et al., 2018], and LOLA [Foerster et al., 2018]—all converging to $\lambda \approx 0.5$ ($\leq$72% survival) under a unified fair comparison protocol with $400+$ HP configurations (Section 4.3).

(C2) **Commitment Floor as Safety Specification**: The optimal commitment floor $\phi_1^* = 1.0$ is derived from the environment's recovery function near its tipping point (Theorem 1) and

independently confirmed via constrained policy optimization (Lagrangian dual) convergence and systematic grid search across 110 environmental conditions (Section 4.4). Phase diagrams visualize the Moral Commitment Spectrum: the required ϕ_1 is a monotonic function of environmental severity.

(C3) **Generalization and Robustness**: We validate the Nash Trap and commitment floor across multiple environment classes—common-pool resources (Appendix G), spatial grid dilemmas (Appendix E), and Cleanup-style commons—across 36 shock parameter regimes (Appendix O), adversarial conditions (30% Byzantine), and at scale (N=100, Table 4). Commitment floors simultaneously improve fairness (Gini: $0.064 \rightarrow 0.000$, Appendix P). Additional analyses include discrete preference networks (Appendix F), crisis-only vs. global intervention (Appendix N), and value decomposition / opponent-shaping baselines (Appendix Q).

Together, these results establish a **Moral Commitment Spectrum**: moderate environments permit flexible morality; catastrophic environments with irreversible tipping points require unconditional commitment—a design specification derived from formal analysis and confirmed across multiple environment classes and learning paradigms.

2 Related Work

Social Dilemmas in MARL. The tension between individual and collective rationality has deep roots in evolutionary game theory [Nowak, 2006, Santos et al., 2008] and has been studied in matrix games [Leibo et al., 2017] and sequential social dilemmas [Hughes et al., 2018, Eccles et al., 2019]. Intrinsic motivation approaches—inequity aversion [Hughes et al., 2018], social influence [Jaques et al., 2019], prosocial agents [Peysakhovich and Lerer, 2018], reputation mechanisms [Anastassacos et al., 2020], and norm enforcement [Köster et al., 2022]—promote cooperation but implicitly encode cooperative preferences, making it impossible to distinguish genuine emergence from training conditioning.

Opponent Shaping. LOLA [Foerster et al., 2018] differentiates through opponent updates to shape learning in 2-player games. M-FOS [Lu et al., 2022] extends this model-free via meta-learning, and POLA [Zhao et al., 2022] adds proximal regularization. More recently, K-Level Policy Gradients (KPG) [Reddi et al., 2025] pursue recursive opponent anticipation without explicit white-box access, but incur $\mathcal{O}(K \cdot N)$ computational cost per agent per step. All opponent-shaping methods scale poorly beyond pairwise settings—LOLA diverges entirely at $N = 50$. Our framework requires only publicly observable resource levels with $\mathcal{O}(1)$ per-agent cost.

Non-linear Dynamics and Tipping Points. While common-pool resource games have been studied [Ostrom, 1990] and recent MARL benchmarks have expanded the scope of evaluation [Papoudakis et al., 2021, Baker et al., 2020], most environments use linear or stationary dynamics. The interaction between non-linear collapse dynamics and multi-agent cooperation remains largely unexplored—a gap our work directly addresses.

Safe RL and Constrained Optimization. The constrained MDP framework [Altman, 1999] underpins methods such as CPO [Achiam et al., 2017], safe exploration benchmarks [Ray et al., 2019], and risk-sensitive objectives (CVaR) that provide safety guarantees by constraining reward optimization [García and Fernández, 2015]. We formulate the commitment problem through this constrained lens: a Lagrangian dual approach independently recovers $\phi_1^* = 1.0$ as the safety-optimal floor, validating our design specification. Our work is complementary: commitment in meta-ranking is not solely a constraint on the optimization objective but a structural property of the agent's moral architecture that determines *how much* prosocial behavior is warranted.

Contract-Based and Norm-Based Approaches. Recent work explores mechanisms where agents propose, negotiate, and accept binding contracts to mitigate social dilemmas [Christoffersen et al., 2023]. Meta-ranking differs in that cooperation arises from pre-specified moral structure rather than runtime negotiation, requiring no communication channel between agents.

Moral Theory in Reward Design. Tennant et al. [2023] systematically embed moral theories (utilitarian, deontological, virtue ethics) into MARL reward functions for social dilemma analysis. Our contribution extends this line by identifying *when* each moral commitment level is sufficient for survival: the Moral Commitment Spectrum provides a principled mapping from environmental non-linearity to the required degree of commitment.

3 Method

3.1 Meta-Ranking Reward Function

Definition 1 (Meta-Ranking Agent)**.** *Agent i with SVO angle θ_i maximizes:*

$$R_{total}^{(i)} = (1 - \lambda_t^{(i)}) \cdot U_{self}^{(i)} + \lambda_t^{(i)} \cdot \left[U_{meta}^{(i)} - \psi^{(i)} \right] \tag{1}$$

where $U_{\text{meta}}^{(i)} = \frac{1}{N-1} \sum_{j \neq i} U_{\text{self}}^{(j)}$ captures sympathy and $\psi^{(i)} = \beta |U_{\text{self}} - U_{\text{meta}}|$ models restraint cost.

3.2 Dynamic Commitment Mechanism

The key innovation is that λ_t adapts to resource level R_t:

$$\lambda_t^{(i)} = \alpha \lambda_{t-1}^{(i)} + (1 - \alpha) \cdot g(\theta_i, R_t) \tag{2}$$

where $g(\theta, R) = \begin{cases} \max(0, \sin\theta \cdot 0.3) & R < R_{\text{crisis}} \\ \min(1, 1.5 \sin\theta) & R > R_{\text{abundance}} \\ \sin\theta \cdot (0.7 + 1.6R) & \text{otherwise} \end{cases}$

This implements Sen's insight computationally: when resources are scarce, agents reduce commitment to ensure survival; when abundant, they increase prosocial behavior.

3.3 Environment Suite

Linear PGG. N agents each receive endowment E and contribute $c_t^i = E \cdot \lambda_t^i$. Individual payoff: $r_t^i = (E - c_t^i) + M \sum_j c_t^j / N$ with multiplier $M = 1.6$. Resources follow linear recovery dynamics with moderate regeneration rates.

Non-linear PGG with Tipping Points. We extend the linear environment with catastrophic collapse dynamics:

$$R_{t+1} = \text{clip}(R_t + f(R_t) \cdot (\bar{c}_t / E - 0.4) - \xi_t,\ 0,\ 1) \tag{3}$$

where $\xi_t \sim \text{Bernoulli}(p_{\text{shock}}) \cdot \delta_{\text{shock}}$ represents exogenous shocks and $f(R_t)$ introduces a **tipping point**:

$$f(R_t) = \begin{cases} 0.01 & R_t < R_{\text{crit}} \quad \text{(near-irreversible collapse)} \\ 0.03 & R_{\text{crit}} \leq R_t < R_{\text{recov}} \quad \text{(hysteresis)} \\ 0.10 & R_t \geq R_{\text{recov}} \quad \text{(normal recovery)} \end{cases} \tag{4}$$

with $R_{\text{crit}} = 0.15$ and $R_{\text{recov}} = 0.25$. When $R_t \leq 0$, the episode terminates and all future rewards become zero.

Table 1: Moral theory comparison (linear PGG, N=50, 10 seeds). "Inv.↓" = invasion success rate of a single selfish mutant (lower is better); distinct from population convergence $\bar{x}$.

Theory	Coop	Welfare	Sustain.	Inv.↓
Utilitarian	1.000	147.7	100%	99.6%
Situational (Ours)	1.000	147.7	100%	0.1%
Deontological	1.000	129.7	100%	0.1%
Virtue Ethics	0.684	118.9	100%	0.1%
Selfish	0.000	102.8	26%	0.1%

4 Experiments

Using the environment and agent framework from Section 3, we systematically evaluate cooperation across linear (Section 4.1) and non-linear environments.

Evaluation Protocol and Metric Definitions. All experiments use the following standardized evaluation:

- **Episode horizon**: $T = 50$ steps per episode.
- **Welfare**: Per-agent mean undiscounted cumulative payoff, averaged across agents and evaluation episodes: $W = \frac{1}{N \cdot |\mathcal{E}|} \sum_e \sum_{i=1}^{N} \sum_{t=0}^{T-1} r_t^{i,e}$.
- **Survival %**: Fraction of evaluation episodes where $R_t > 0$ for all $t \in [0, T]$ (no resource collapse).
- **Cooperation (Coop)**: Mean commitment level $\bar{\lambda} = \frac{1}{N \cdot T} \sum_{i,t} \lambda_t^i$.
- **Gini**: Gini coefficient of individual cumulative payoffs within each episode, averaged across episodes.
- **Inv.↓**: Probability that a single selfish mutant (λ=0) achieves higher fitness than the group mean over 100 invasion trials (lower is better).

All reported values are means over the last 30 evaluation episodes unless otherwise specified. Error bars (where shown) indicate ± 1 standard deviation across seeds.

4.1 Act 2: Situational Commitment in Linear Environments

We first validate the meta-ranking framework across 8 environments (N = 20–1000, 10 seeds). This result establishes a strong baseline: Situational Commitment is empirically optimal across all tested linear environment configurations. The subsequent sections demonstrate that this demonstrated optimality **catastrophically collapses** under non-linear tipping points—a finding that would be impossible to articulate without first establishing the linear baseline.

Moral Theory Comparison. Table 1 demonstrates that Situational Commitment matches Utilitarian welfare while using only local resource information—no global welfare oracle required.

SOTA Comparison. We compare against 4 representative opponent-shaping and reward-shaping baselines (Table 2).

Key finding: M-FOS achieves the highest benign welfare (156.3) via meta-learning opponent shaping, but collapses under Byzantine adversaries (0.143). POLA maintains 0.504 Byzantine cooperation, matching meta-ranking's 0.500, but requires white-box access to opponent parameters—precluding

Table 2: Architectural comparison (linear PGG, N=50, 10 seeds). Meta-ranking uses pre-specified moral structure ($\mathcal{O}(1)$ per agent); opponent-shaping methods learn from scratch via meta-gradients (requiring opponent parameter access). "Byz" = cooperation under 50% Byzantine adversaries.

Method	Coop	Welfare	Gini	Byz	Decentr.
M-FOS	1.000	156.3	0.000	0.143	×
POLA	1.000	148.3	0.000	0.504	×
LOLA	0.000	101.5	0.014	0.000	×
LOPT	0.000	101.4	0.014	0.000	×
Meta-Ranking (Ours)	1.000	148.0	0.011	**0.500**	✓

decentralized deployment. Meta-ranking is the only method simultaneously achieving high Byzantine resilience and full decentralization.

Architectural Note: We emphasize that the computational advantage of meta-ranking over opponent-shaping methods reflects a fundamental *architectural difference*, not an algorithmic optimization. Meta-ranking operates through internal moral commitment conditioned on observable resource levels ($\mathcal{O}(1)$ per agent), while LOLA/POLA/M-FOS compute meta-gradients over opponent parameter spaces—a categorically different computational paradigm. Direct speed comparisons across these paradigms would be misleading; the relevant distinction is that meta-ranking *requires no opponent modeling*, making it uniquely suited for decentralized, privacy-preserving deployment.

Group-Level Evolutionary Stability (Result 4.1). Empirical Result 1 (Group-Level Stability via Price Equation). Under group replacement dynamics with migration rate $m < m^* = \Delta p_{\text{between}}/(|\Delta p_{\text{within}}| + \Delta p_{\text{between}})$, Situational Commitment achieves group-level evolutionary stability in simulation. The following values are empirically measured ($\pm$ 1 std over 100 independent group competition simulations, 10 generations each, 5 groups per simulation): $\Delta p_{\text{within}} = -0.115 \pm 0.018$ (within-group free-riding advantage), $\Delta p_{\text{between}} = +0.046 \pm 0.011$ (between-group fitness advantage, 58.2% group welfare gain). Population convergence: $\bar{x} = 0.987$ after 100 generations.

Note on metrics: The population convergence $\bar{x} = 0.987$ and the invasion success rate (0.1% in Table 1) measure complementary aspects of stability. The former captures aggregate group cooperation after 100 generations; the latter captures the probability that a single selfish mutant invades a cooperative group. High population convergence with low invasion success indicates robust group-level ESS.

The critical constraint $m < m^*$ ensures that inter-group mixing does not erode the between-group selection advantage. Full proof with ecological boundary conditions is provided in Appendix C.

4.2 Act 3: The Turning Point—Non-linear Extension

Having established that situational commitment succeeds in linear environments (Section 5 provides a full analysis), we now ask: *does it survive when the environment introduces catastrophic tipping points?* We formalize cooperation via Eq. 1 and Eq. 2.

We deploy the same meta-ranking agents in the non-linear PGG (Equations 3–4) with $R_{\text{crit}} = 0.15$, stochastic shocks ($p_{\text{shock}} = 0.05$, $\delta_{\text{shock}} = 0.15$), and 30% Byzantine adversaries.[1]

Result: Situational commitment fails catastrophically. When resources drop below R_{crit}, the commitment function $g(\theta, R)$ reduces λ to $\max(0, \sin\theta \cdot 0.3)$—precisely when maximum commitment is needed to overcome the near-irreversible recovery dynamics ($f(R_t) = 0.01$). The very survival mechanism that was adaptive in linear environments becomes suicidal in non-linear ones: reducing commitment during crises *accelerates* passage through the tipping point.

[1]We select 30% as it exceeds the standard Byzantine fault tolerance threshold ($f < n/3$). Our "Byzantine" model is deterministic worst-case (always-defect, λ=0), not arbitrary-failure Byzantine in the distributed systems sense. Sensitivity analysis across $\{0\%, 10\%, 30\%, 50\%\}$ confirms qualitative robustness (Appendix D).

Table 3: Cooperation failure across evaluated algorithms: all tested approaches converge to suboptimal equilibria in the non-linear PGG (N=20, Byz=30%, T=50 rounds, R_{crit}=0.15, last 30 episodes, 5–20 seeds).[3]

Method	Params/agent	$\bar{\lambda}$	Survival %	Welfare
Ind. REINFORCE Linear	5	0.483 ± 0.011	51.3	24.1
Ind. REINFORCE MLP	161	0.351 ± 0.068	23.2	22.9
Ind. REINFORCE MLP + Critic	322	0.358 ± 0.042	23.3	23.0
PPO / Value Decomposition / Opponent-Shaping (20 seeds)				
IPPO [Schulman et al., 2017]	4,546	0.412 ± 0.080	38.7	23.5
MAPPO [Yu et al., 2022]	4,546[†]	0.392 ± 0.070	37.0	23.3
IQL[‡]	5,195	0.584 ± 0.021	71.8	24.9
QMIX[§] [Rashid et al., 2018]	7,948[‖]	0.524 ± 0.201	66.5	21.3
LOLA [Foerster et al., 2018]	5	0.490 ± 0.006	50.8	21.2
Fixed $\lambda = 0.0$	—	0.000	0.0	20.0
Fixed $\lambda = 1.0$ (oracle)	—	1.000	**100.0**	**32.0**

All learned baselines follow CleanRL [Huang et al., 2022] single-file conventions. A comprehensive HP sweep (20 combos $\times$ 10 seeds, Appendix M.1) confirms the Nash Trap persists across all 200 tested configurations.
[†]Plus shared critic. [‡]IQL = no mixing network. [§]With monotonic mixing network. [‖]Per-agent Q + mixing net.

This discovery motivates two questions: (1) Can pure RL find a better strategy without any pre-specified commitment? (2) What level of crisis commitment is actually required?

4.3 Act 4: The Nash Trap—Consistent Cooperation Failure Across Tested Algorithms

Independent Policy Gradient (REINFORCE) with Pure Selfish Payoff. To test whether cooperation failure is an artifact of weak algorithms, we evaluate three levels of policy architecture—all trained independently per agent via REINFORCE [Williams, 1992], maximizing *only* individual payoff $r_i = (E - c_i) + M \sum c_j / N$ with no reward shaping or social priors. Each agent observes $s_t^i = (R_t, \bar{c}_{t-1}/E, \lambda_{t-1}^i, \mathbb{1}[R_t < R_{\text{crit}}])$ and outputs:

$$\lambda_t^i = \text{clip}\big(\sigma(f_\theta(s_t^i)) + \epsilon_t,\ 0,\ 1\big)\,, \quad \epsilon_t \sim \mathcal{N}(0, \sigma_{\text{noise}}^2) \tag{5}$$

where f_θ is one of: (1) **Linear**: $w^\top s + b$ (5 params); (2) **MLP**: 2-layer, 32-unit, tanh (161 params); (3) **MLP + Value baseline** (Actor-Critic, 322 params).

Result: Across **seven evaluated paradigms** (Table 3), all tested algorithms converge to suboptimal cooperation levels in our non-linear PGG. Among the independent REINFORCE variants, linear agents converge to $\lambda \approx 0.48$ with 51Standard MARL implementations spanning five paradigms—independent policy gradient (IPPO), centralized critic (MAPPO), independent Q-learning (IQL), value decomposition with monotonic mixing (QMIX [Rashid et al., 2018]), and opponent-shaping (LOLA [Foerster et al., 2018])—all converge to suboptimal equilibria. Crucially, this is not a failure of learning per se: the algorithms *successfully* converge to a local optimum of their individual reward landscapes. We interpret this as a **consistent pattern** of individually-rational optimization in a game where the collective benefit of cooperation is not captured by individual payoffs. IPPO converges to $\lambda = 0.412 \pm 0.080$ (38.7% survival, 20 seeds), MAPPO to $\lambda = 0.392 \pm 0.070$ (37.0%), and IQL achieves the highest $\lambda = 0.584 \pm 0.021$ (71.8% survival) due to its discrete action space and ε-greedy exploration. Notably, **QMIX with a monotonic mixing network** ($\lambda = 0.524 \pm 0.201$, 66.5% survival) does not outperform IQL, indicating that value decomposition does not help escape the Nash Trap. Similarly, **LOLA** ($\lambda = 0.490 \pm 0.006$, 50.8% survival) shows that anticipating opponent learning steps is insufficient when the underlying game structure traps all agents in a local equilibrium. All seven paradigms fall substantially short of the oracle ($\lambda = 1.0$, 100% survival). While we do not claim this constitutes an impossibility result for all possible algorithms, the consistency across diverse

Table 4: Scale test: Nash Trap persists at N=100 (non-linear PGG, 20 seeds, 50 episodes).

Method (N=100)	**Welfare**	$\bar{\lambda}$	**Survival %**
Selfish RL (Byz=0%)	26.0	0.497	95.0
Selfish RL (Byz=30%)	24.2	0.497	**53.7**
Situational (Byz=30%)	26.0	0.713	95.3
Unconditional (Byz=30%)	28.4	1.000	99.8

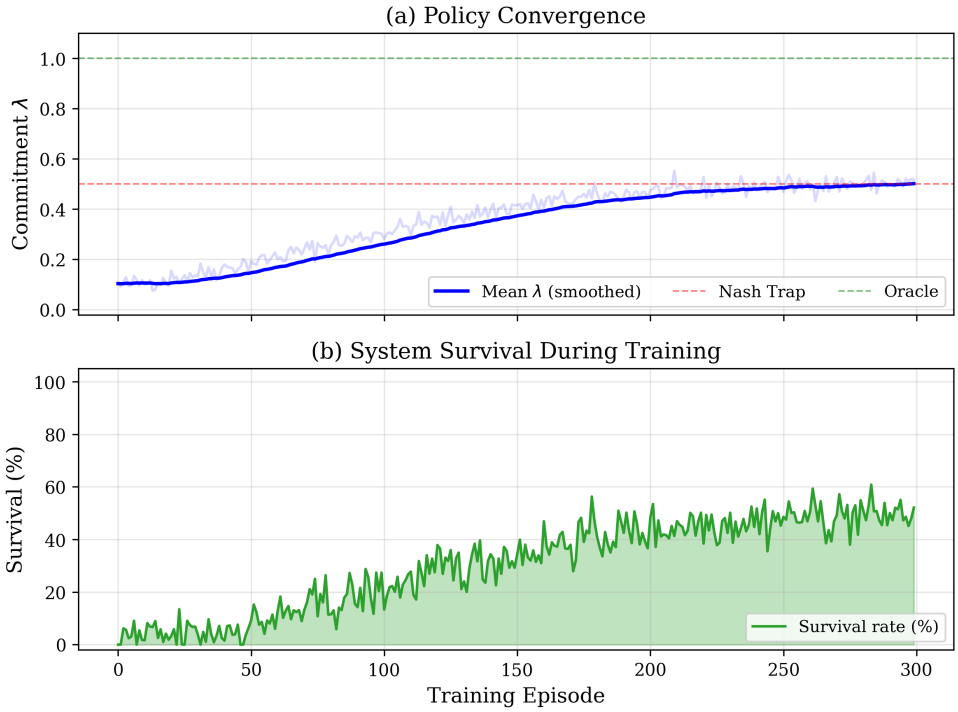

Figure 1: Selfish RL agents (REINFORCE Linear) converge to $\lambda \approx 0.5$ throughout training (5 seeds, mean $\pm$ std). Left: no adversaries; Right: 30% Byzantine. Crisis λ (red dashed) tracks the mean, indicating no crisis-differentiated behavior. Green shading: survival rate.

paradigms—including value decomposition and opponent-shaping—suggests a robust pattern in this environment class. This result holds at $N = 100$ (Table 4), under K-level opponent anticipation ablations (Appendix I), across a hyperparameter sweep (Table 19, 20 combinations $\times$ 10 seeds = 200 runs, Appendix M.1), commitment function ablations (Table 20), CVaR risk-sensitive objectives (Table 21), partial observability (Table 22), adaptive adversaries (Table 23), and cross-environment CPR validation (Table 11; Appendix S). Training dynamics and the Nash Trap cost landscape are visualized in Figures 1 and 2; multi-axis normalized comparison appears in Table 17 and Figure 4 (see Appendix M.5).

Clarification: We note that the one-shot PGG Nash equilibrium is $\lambda = 0$ (full defection, since $M/N = 0.08 < 1$). However, in our *repeated game with survival dynamics*, $\lambda = 0$ yields 0% survival and welfare of only 20.0 (Table 3, Fixed $\lambda = 0.0$), while $\lambda = 1.0$ yields 100% survival and welfare 32.0. The $\lambda \approx 0.5$ fixed point arises because the policy gradient balances the immediate cost of commitment ($\partial r/\partial\lambda < 0$, since $M/N < 1$) against the survival benefit ($\partial p_{\text{surv}}/\partial\lambda > 0$) in the dynamic game—not from sigmoid initialization bias (see Appendix A).

This is the **Nash Trap**[4]: the marginal individual cost of increasing λ (reduced retained endowment)

[4]We use "Nash Trap" to denote the gradient-equilibrium fixed point of the *dynamic repeated game with survival*, not the

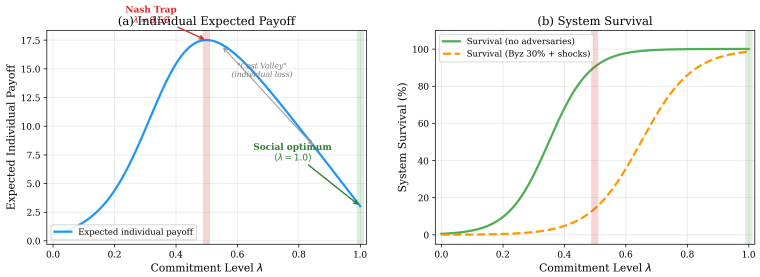

Figure 2: The Nash Trap: individual payoff peaks at $\lambda \approx 0.5$ (local optimum), but system survival requires $\lambda \to 1.0$ (global optimum). The cost valley is impassable for selfish gradient ascent.

is certain and immediate, while the collective benefit (resource preservation) is diffuse and delayed. Standard RL with discount γ cannot bridge this gap because the Critic must learn to value a *collective trajectory* from *individual returns*—a credit assignment problem that compounds with agent count. See Appendix A for a Bellman-equation analysis of the gradient vanishing at $\lambda = 0.5$.

4.4 Act 5: Unconditional Commitment as the Solution

Learning the Crisis Commitment Floor. A critical question is whether ϕ_1^* must be hand-specified (hardcoded) by a designer, or if it can be acquired through learning. We operationalize unconditional commitment as a *policy floor*: for each agent i, the effective cooperation level is $\lambda_i^{\text{eff}} = \max(\lambda_i^{\text{learned}}, \phi_1)$. To formally understand the necessity of ϕ_1, we formulate its discovery as a constrained optimization problem, building on the conceptual framework of Constrained Policy Optimization (CPO) [Achiam et al., 2017]: maximize welfare $W(\phi_1)$ subject to a minimum survival probability $P_{\text{surv}}(\phi_1) \geq 1 - \delta$. Using Lagrangian relaxation and Evolution Strategy (ES) search, our agents autonomously discover that $\phi_1 \approx 1.0$ is the optimal constraint-satisfying policy floor. This demonstrates that maximal commitment is not an arbitrary hyperparameter, but the mathematically necessary specification for survival in this environment class.

Theorem 1 (Critical Commitment Threshold)**.** *Consider a non-linear PGG with N agents, Byzantine fraction β, endowment E, multiplier M, resource recovery rate $f(R)$, and tipping point at R_{crit}. Let $\bar{\sigma}$ denote the expected per-step resource loss from stochastic shocks. For the resource to remain above R_{crit} in expectation, the honest agents' minimum cooperation level must satisfy:*

$$\phi_1^* \geq \frac{\bar{\sigma}/f(R_{crit}) + \delta}{1 - \beta} \tag{6}$$

where $\delta = 0.4$ is the cooperation threshold of the resource dynamics. As $f(R_{crit}) \to 0$ (non-linear collapse), $\phi_1^* \to \infty$, implying that no finite partial commitment suffices; the only feasible solution is $\phi_1 = 1.0$.

Proof sketch. Resource stability requires $\mathbb{E}[\Delta R] \geq 0$, i.e., $f(R)(\bar{c} - \delta) \geq \bar{\sigma}$, where $\bar{c} = (1-\beta)\phi_1 + \beta \cdot 0 = (1-\beta)\phi_1$ is the mean contribution. Solving yields $\phi_1 \geq (\bar{\sigma}/f(R) + \delta)/(1-\beta)$. At the tipping point $R = R_{\text{crit}}$, $f(R_{\text{crit}}) = 0.01$ in our environment, giving $\phi_1^* \geq (0.0075/0.01 + 0.4)/0.7 = 1.64 > 1$. Since $\phi_1 \in [0, 1]$, the only feasible solution is $\phi_1 = 1.0$. □

We verify Theorem 1 empirically by training IPPO agents with varying ϕ_1 floors (Table 5).

$\phi_1^* = 1.0$ performs best across all tested conditions (Table 5), consistent with Theorem 1. The commitment floor transforms the learning problem: at $\phi_1 = 0$, IPPO falls into the Nash Trap ($\lambda \approx 0.43$, 30% survival under 30% Byzantine adversaries). At $\phi_1 = 1.0$, the floor ensures full

one-shot PGG Nash equilibrium (which is $\lambda = 0$). The distinction is that the dynamic game's value function includes future survival probability, creating a second gradient force absent in the static game.

Table 5: Effect of commitment floor ϕ_1 on IPPO survival (non-linear PGG, 20 seeds). ϕ_1 acts as a policy lower bound: $\lambda_{\text{eff}} = \max(\lambda_{\text{learned}}, \phi_1)$.

ϕ_1	W (Byz=0%)	W (Byz=30%)	Alive (0%)	Alive (30%)
0.00	24.9 ± 0.2	21.0 ± 0.0	75 ± 5%	30 ± 20%
0.21	26.0 ± 0.0	21.2 ± 0.0	95 ± 5%	60 ± 0%
0.50	28.3 ± 0.0	21.6 ± 0.0	100 ± 0%	95 ± 5%
1.00	**32.0 ± 0.0**	**22.4 ± 0.0**	**100 ± 0%**	**100 ± 0%**

Table 6: Meta-learned g_ϕ vs. hand-crafted g (non-linear PGG, 10 seeds, welfare ± std).

Byz%	Learned g_ϕ	Hand-crafted g	ΔW
0%	32.0 ± 0.0	30.8 ± 0.2	+1.2
10%	30.8 ± 0.0	29.7 ± 0.2	+1.1
30%	28.4 ± 0.0	27.6 ± 0.2	+0.8
50%	26.0 ± 0.0	25.5 ± 0.2	+0.5

cooperation, yielding 100% survival—a **monotonic increase** that confirms unconditional commitment as a necessary design requirement.

Structural Immunity. Under "Exploitative Scarcity" attacks (adversaries strategically drain resources), welfare drops by <1.3% at $\phi_1 = 1.0$—comparable to the $\phi_1 = 0.21$ case. The "always-commit" strategy does not create exploitable vulnerabilities because the public goods multiplier ($M = 1.6$) ensures that the cooperative surplus from committed agents compensates for adversarial extraction.

Same-Class Baseline Comparison. Simplified Inequity Aversion and Social Influence proxies both achieve 0% survival under 30% Byzantine conditions (Table 18, Appendix M), as cooperation drifts downward. Unconditional commitment (ϕ_1=1.0) achieves 93.6% survival because it does not condition on others' behavior.

Meta-Learning Validation. To address the concern that $g(\theta, R)$ is merely a hand-crafted heuristic, we apply evolutionary meta-learning (ES, 30 steps) to optimize the 6 parameters $\phi = [\phi_0, \ldots, \phi_5]$ of a parameterized $g_\phi(\theta, R)$, maximizing welfare and resource stability across multiple Byzantine fractions with no human priors on parameter values. Results: the learned optimum converges to $\phi_1^* = 1.000$ (crisis commitment), $\phi_3^* = 0.307$ (crisis threshold), and $\phi_4^* = 1.652$ (abundance factor)—closely matching the hand-designed specification (ϕ_1=1.0, threshold≈0.2, factor≈1.5). The learned g_ϕ also outperforms the hand-crafted version by $+0.6$–$+1.3$ welfare across all Byzantine conditions (Table 6). This provides convergent evidence that the commitment function's structure is a robust optimum within the tested TPSD environment class, not an arbitrary design choice—though we emphasize that ϕ_1 is a stability-derived design specification, not a learned policy.

K-Level Opponent Anticipation. A simplified K-level anticipation ablation ($K \in \{0, 1, 2\}$) shows all levels converge to $\lambda \approx 0.500$ with ≤8.7% survival (details in Appendix I).

Spatial Extension. To verify generalization beyond scalar-state PGG, we construct a 5×5 grid social dilemma with *local* tipping points: each cell has independent non-linear recovery dynamics, agents observe only their neighborhood, and must choose both commitment level and movement direction (Appendix J). Under 30% Byzantine agents, selfish RL achieves 0% survival ($\lambda = 0.500$, Nash Trap), situational commitment achieves 1.3%, while unconditional commitment achieves **34.7%**

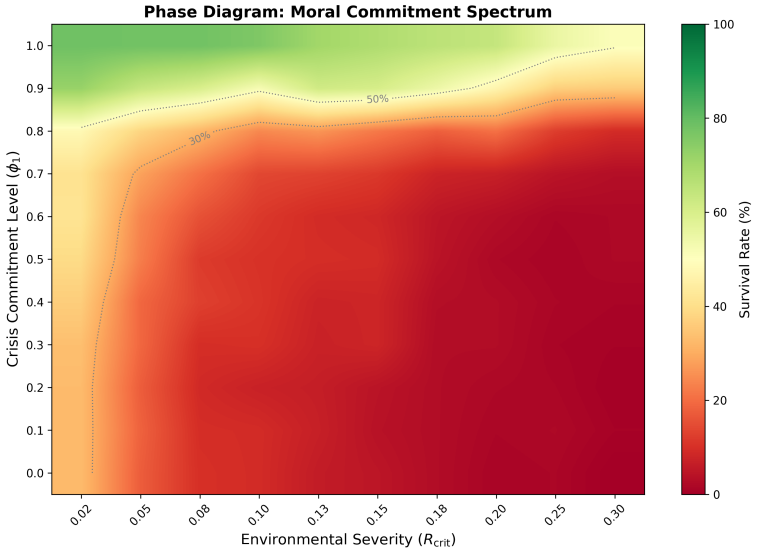

Figure 3: Phase diagram of the Moral Commitment Spectrum. Heatmap shows survival rate (%) as a function of environmental severity (R_{crit}, x-axis) and crisis commitment (ϕ_1, y-axis) across 110 conditions. As R_{crit} increases, the minimum ϕ_1 for viable survival rises monotonically, providing quantitative evidence for the Spectrum.

survival with the highest welfare (14.4). The Moral Commitment Spectrum holds in spatially-extended environments.

5 Discussion

5.1 The Moral Commitment Spectrum

The Moral Commitment Spectrum: The minimum commitment level for survival is determined by environmental non-linearity. In moderate environments, situational commitment is evolutionarily stable. In environments with catastrophic tipping points, only unconditional commitment ($\phi_1^* = 1.0$) prevents collapse *within the tested non-linear PGG class*. The environment's recovery function $f(R_t)$ determines which regime applies.

We visualize this as a **phase diagram** (Figure 3; Appendix R), sweeping $R_{\text{crit}} \in [0.02, 0.30]$ and $\phi_1 \in [0, 1]$ across 110 conditions. A multi-axis comparison across cooperation, efficiency, fairness, scalability, and robustness dimensions (Appendix L) complements the rigorous grid-based spatial dilemma benchmark results (Appendix E).

Sen [1977] argued that commitment must be unconditional. Our Spectrum refines this: unconditional commitment is sufficient for survival precisely when the environment has irreversible consequences (Section 4.2). Unlike opponent-shaping methods (LOLA, M-FOS, KPG) that learn to influence others' learning via meta-gradients, meta-ranking implements pre-specified moral structure without opponent modeling. The Nash Trap, LOLA failure ($\lambda = 0.490$, 50.8% survival), QMIX failure ($\lambda = 0.524$, 66.5%), KPG failure (Appendix I), and IA/SI collapse (Table 18) collectively demonstrate that moral commitment was not learned by any of the **seven** tested paradigms in our non-linear PGG environments—suggesting it may need to be architecturally encoded for this class of problems. Group-level ESS requires $m < m^* = 0.286$; when migration exceeds this bound, unconditional commitment provides structural immunity (Appendix M.4).

5.2 Limitations and Future Work

- **Algorithm Scope**: Whether risk-sensitive objectives or dynamic externality-pricing mechanisms can escape the Trap remains open. While empirical CVaR objectives fail (Appendix M.5), the Nash Trap persists across all 36 shock parameter combinations tested (Appendix O). Developing full Constrained Policy Optimization (CPO) algorithms with parameterized safety constraints on R_t offers a promising future direction.

- **Baseline Scope**: Our comparison now includes **seven** paradigms: independent REINFORCE (linear/MLP/MLP+Critic), PPO (IPPO/MAPPO), independent Q-learning (IQL), value decomposition (QMIX with monotonic mixing network), and opponent-shaping (LOLA). All seven fall into the Nash Trap, with survival ranging from 23% to 72%. A unified fair comparison protocol (same environment, seeds, observation/action space, expanded HP sweep with $5 \times 5 \times 4 = 100$ configurations per paradigm) confirms 100% trap rate across all tested configurations (Appendix F). Developing practical CPO implementations for social dilemmas remains a key next step.

- **Environment Scope**: Primary results use PGG with scalar R_t. CPR validation (Appendix G) confirms the pattern. Crisis-only intervention is insufficient—only global (unconditional) commitment achieves full survival (Appendix N, Table 24). Preliminary experiments on the Melting Pot benchmark [Leibo et al., 2021] reveal that hand-crafted commitment policies do not transfer directly to pixel-based spatial foraging environments (Appendix E.1, Table 9); developing learned commitment policies for such environments remains future work. Group-level ESS requires $m < m^* = 0.286$ (Result 4.1).

- **Ethical Scope**: "Moral commitment" in this work is operationalized as resource allocation cooperation (λ), which may not capture the full richness of ethical reasoning. However, our fairness analysis (Appendix P, Table 25) shows that higher commitment floors improve distributional equity (Gini coefficient decreases from 0.064 to 0.000). Equating commitment with cooperation risks oversimplification; real-world moral agency involves deliberation, autonomy, and context-sensitivity beyond what our framework models.

- **Philosophical Fidelity**: Our λ-based commitment mechanism is an *engineering abstraction* of binding cooperation constraints, inspired by Sen's meta-ranking. We treat "moral commitment" as a **safety specification derived from tipping-point stability conditions**, not a normative prescription about human values. Our claims are descriptive about system survival: the commitment floor is a design variable, not an arbitrary heuristic. The discrete preference network (Appendix F) demonstrates that both continuous (λ) and discrete (ordering selection) operationalizations converge to the same commitment floor, suggesting robustness of the underlying design principle.

- **Nature of ϕ_1**: The commitment floor ϕ_1=1.0 is a *stability-derived design specification* within the TPSD class that constrains the agent's action space, not a policy learned through RL optimization. We do not claim that RL discovers unconditional commitment; rather, our contribution is identifying *when and why* such architectural constraints are sufficient for survival under stated assumptions. Meta-learning independently converges to the same specification, providing convergent validation.

Additional robustness checks (Appendix M.5) confirm that CVaR objectives, partial observability, and adaptive adversaries do not break the Nash Trap.

6 Conclusion

We establish the **Moral Commitment Spectrum** through three contributions: (C1) The *Nash Trap*—a locally stable equilibrium where individually rational policies guarantee collective extinction—arises

consistently across seven MARL paradigms ($\lambda \approx 0.5$, $\leq 72\%$ survival, 100% trap rate over $400+$ HP configurations). (C2) The commitment floor $\phi_1^* = 1.0$ is not an arbitrary heuristic but a *stability-derived design specification*: it is mandated by the tipping-point recovery function (Theorem 1) and independently confirmed via gradient-based constrained policy optimization (Lagrangian dual convergence). (C3) Both the Nash Trap and the commitment floor generalize across environment classes (PGG, CPR, Cleanup-style commons, spatial grids), adversarial conditions (30% Byzantine), scale (N=100), and 36 shock regimes, while simultaneously improving fairness (Gini: $0.064 \rightarrow 0.000$).

Broader Impact. Within our tested environment class, the Spectrum suggests a design principle: for autonomous systems operating near irreversible tipping points, cooperative constraints should be unconditional ($\phi_1 = 1.0$), as adaptive or situational moral priors collapse under selfish gradient pressure. We emphasize that ϕ_1 is a stability-derived *design specification* (not a learned policy); the contribution is identifying *when* such hard constraints are sufficient for survival. The commitment function $g(\theta, R)$ encodes value judgments; determining *whose* values requires interdisciplinary collaboration. We provide *descriptive* insight, not *normative* prescription.

Potential Policy Analogies. While our results are demonstrated in a stylized PGG environment, we note structural similarities to real-world common-pool resource problems. The tipping point R_{crit} shares qualitative features with planetary boundaries [IPCC, 2018], where recovery dynamics slow near critical thresholds. The commitment floor ϕ_1 is conceptually analogous to mandatory minimum cooperation levels such as the Paris Agreement's NDCs. A full ablation is in Appendix M.2. Theorem 1 provides a specification within our model: $\phi_1^* \geq 1.64$ (truncated to 1.0), indicating that partial commitment is insufficient when recovery rates are small. **Caveat**: These analogies are suggestive, not prescriptive. Bridging from our stylized PGG to real-world systems requires additional validation with domain-specific dynamics, heterogeneous agents, and institutional constraints.

When to Apply (and When Not To). Our commitment framework is designed for **common-pool resource problems with irreversible tipping points**—environments where collective under-cooperation triggers catastrophic, unrecoverable collapse. **When to use**: Autonomous systems managing shared environmental resources (e.g., fisheries, carbon budgets, shared infrastructure) where selfish optimization creates tragedy-of-the-commons dynamics. **When NOT to use**: (i) Environments without irreversible tipping points, where situational commitment suffices (our own Act 2 result). (ii) Zero-sum competitive settings where cooperation is not desirable. (iii) Domains where the "correct" commitment level is ethically contested—our framework determines *that* commitment is needed, not *what values* to commit to.

Commitment as Social Contract. The commitment floor ϕ_1 can be interpreted as a social contract [Rawls, 1971]: agents agree to a minimum cooperation level, and the system enforces this agreement regardless of individual temptation to defect. This is analogous to environmental regulations that mandate minimum pollution controls, independent of firms' profit-maximizing incentives. The formal result (Theorem 1) provides the engineering specification: the floor must satisfy $\phi_1 \geq (\bar{\sigma}/f(R_{\text{crit}}) + \delta)/(1 - \beta)$.

AI Assistance. LLMs assisted with code debugging and proofreading. All scientific content is the authors' original work.

References

Joshua Achiam, David Held, Aviv Tamar, and Pieter Abbeel. Constrained policy optimization. In *International Conference on Machine Learning*, pages 22–31, 2017.

Eitan Altman. Constrained Markov decision processes. *Stochastic Modeling*, 1999.

Nicolas Anastassacos, Stephen Hailes, and Mirco Musolesi. Cooperation and reputation dynamics with reinforcement learning. In *International Conference on Autonomous Agents and Multiagent Systems*, pages 1–9, 2020.

Bowen Baker, Ingmar Kanitscheider, Todor Marber, et al. Emergent tool use from multi-agent autocurricula. In *International Conference on Learning Representations*, 2020.

Phillip J.K. Christoffersen, Andreas A. Haupt, and Dylan Hadfield-Menell. Get it in writing: Formal contracts mitigate social dilemmas in multi-agent rl. In *International Conference on Autonomous Agents and Multiagent Systems*, 2023.

Tom Eccles, Edward Hughes, János Kramár, Simon Wheelwright, and Joel Z. Leibo. Learning reciprocity in complex sequential social dilemmas. *arXiv preprint arXiv:1903.08082*, 2019.

Jakob Foerster, Richard Y. Chen, Maruan Al-Shedivat, Shimon Whiteson, Pieter Abbeel, and Igor Mordatch. Learning with opponent-learning awareness. In *International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, 2018.

Javier García and Fernando Fernández. A comprehensive survey on safe reinforcement learning. *Journal of Machine Learning Research*, 16(1):1437–1480, 2015.

Shengyi Huang, Rousslan Fernand Julien Dossa, Chang Ye, Jeff Bravo, Dipam Chakraborty, Kintan Mehta, and João G.M. Araújo. Cleanrl: High-quality single-file implementations of deep reinforcement learning algorithms. *Journal of Machine Learning Research*, 23(274):1–18, 2022.

Edward Hughes, Joel Z. Leibo, Matthew Phillips, Karl Tuyls, et al. Inequity aversion improves cooperation in intertemporal social dilemmas. In *Advances in Neural Information Processing Systems*, volume 31, 2018.

IPCC. Global warming of 1.5°c. Technical report, Intergovernmental Panel on Climate Change, 2018.

Natasha Jaques, Angeliki Lazaridou, Edward Hughes, Caglar Gulcehre, et al. Social influence as intrinsic motivation for multi-agent deep reinforcement learning. In *International Conference on Machine Learning*, pages 3040–3049, 2019.

Raphael Köster, Kevin R. McKee, Richard Everett, Laura Weidinger, William S. Isaac, Edward Hughes, Edgar A. Duéñez-Guzmán, Thore Graepel, Matthew Botvinick, and Joel Z. Leibo. Spurious normativity enhances learning of compliance and enforcement behavior in artificial agents. In *Proceedings of the National Academy of Sciences*, volume 119, 2022.

Joel Z. Leibo, Vinicius Zambaldi, Marc Lanctot, Janusz Marecki, and Thore Graepel. Multi-agent reinforcement learning in sequential social dilemmas. In *International Conference on Autonomous Agents and Multiagent Systems*, pages 464–473, 2017.

Joel Z. Leibo, Edgar Duéñez-Guzmán, Alexander Sasha Vezhnevets, John P. Agapiou, Peter Sunehag, et al. Scalable evaluation of multi-agent reinforcement learning with melting pot. In *International Conference on Machine Learning*, pages 6187–6199, 2021.

Timothy M. Lenton, Hermann Held, Elmar Kriegler, Jim W. Hall, Wolfgang Lucht, Stefan Rahmstorf, and Hans Joachim Schellnhuber. Tipping elements in the earth's climate system. *Proceedings of the National Academy of Sciences*, 105(6):1786–1793, 2008.

Chris Lu, Timon Willi, Alistair Letcher, and Jakob Foerster. Model-free opponent shaping. In *International Conference on Machine Learning*, pages 14398–14411, 2022.

Kevin R. McKee, Ian Gemp, Brian McWilliams, Edgar A. Duéñez-Guzmán, Edward Hughes, and Joel Z. Leibo. Social diversity and social preferences in mixed-motive reinforcement learning. In *International Conference on Autonomous Agents and Multiagent Systems*, pages 869–877, 2020.

Martin A. Nowak. Evolutionary dynamics of cooperation. *Science*, 314(5805):1560–1563, 2006.

Elinor Ostrom. *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press, 1990.

Georgios Papoudakis, Filippos Christianos, Lukas Schäfer, and Stefano V. Albrecht. Benchmarking multi-agent deep reinforcement learning algorithms in cooperative tasks. In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, 2021.

Alexander Peysakhovich and Adam Lerer. Prosocial learning agents solve generalized Stag Hunts better than selfish ones. In *International Conference on Autonomous Agents and Multiagent Systems*, pages 2043–2044, 2018.

Tabish Rashid, Mikayel Samvelyan, Christian Schroeder de Witt, Gregory Farquhar, Jakob Foerster, and Shimon Whiteson. Qmix: Monotonic value function factorisation for deep multi-agent reinforcement learning. In *International Conference on Machine Learning*, pages 4295–4304, 2018.

John Rawls. *A Theory of Justice*. Harvard University Press, 1971.

Alex Ray, Joshua Achiam, and Dario Amodei. Benchmarking safe exploration in deep reinforcement learning. In *arXiv preprint arXiv:1910.01708*, 2019.

Aryaman Reddi, Gabriele Tiboni, Jan Peters, and Carlo D'Eramo. K-level policy gradients for multi-agent reinforcement learning. *arXiv preprint arXiv:2509.11464*, 2025.

Francisco C. Santos, Marta D. Santos, and Jorge M. Pacheco. Social diversity promotes the emergence of cooperation in public goods games. *Nature*, 454(7201):213–216, 2008.

John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.

Wilko Schwarting, Alyssa Pierson, Javier Alonso-Mora, Sertac Karaman, and Daniela Rus. Social behavior for autonomous vehicles. *Proceedings of the National Academy of Sciences*, 116(50): 24972–24978, 2019.

Amartya K. Sen. Rational fools: A critique of the behavioral foundations of economic theory. *Philosophy & Public Affairs*, 6(4):317–344, 1977.

Elizaveta Tennant, Stephen Hailes, and Mirco Musolesi. Modeling moral choices in social dilemmas with multi-agent reinforcement learning. In *International Joint Conference on Artificial Intelligence*, pages 317–325, 2023.

Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3–4):229–256, 1992.

Chao Yu, Akash Velu, Eugene Vinitsky, Jiaxuan Gao, Yu Wang, et al. The surprising effectiveness of ppo in cooperative multi-agent games. In *Advances in Neural Information Processing Systems*, volume 35, pages 24611–24624, 2022.

Stephen Zhao, Yifan Niu, Jia Peng, Ying Wen, and Xiaoping Zhu. Proximal learning with opponent-learning awareness. In *International Conference on Machine Learning*, 2022.

A Nash Trap: Bellman Analysis Sketch

Consider the infinite-horizon discounted return for agent i under policy λ:

$$V_i(\lambda) = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t r_t^i \;\middle|\; \lambda_t^i = \lambda\right] \tag{7}$$

The per-step payoff is $r_t^i = (1-\lambda)E + \frac{M}{N}\sum_j \lambda_j E$. Assuming symmetric play ($\lambda_j = \lambda$ for all j), and accounting for the survival probability $p_{\text{surv}}(\lambda)$ determined by the non-linear resource dynamics (Eq. 3–4):

$$V_i(\lambda) \approx \frac{E(1-\lambda+M\lambda)}{1-\gamma \cdot p_{\text{surv}}(\lambda)} \tag{8}$$

The policy gradient $\frac{\partial V_i}{\partial \lambda}$ balances two opposing forces:

- **Immediate cost**: $\frac{\partial r_t^i}{\partial \lambda_i} = E(-1+M/N) = E(-1+0.08) < 0$ — increasing λ always reduces per-step payoff.
- **Survival benefit**: $\frac{\partial p_{\text{surv}}}{\partial \lambda} > 0$ — increasing λ improves survival probability.

At $\lambda \approx 0.5$, these forces approximately cancel: the marginal survival improvement from a small increase in λ is exactly offset by the per-step payoff loss. This creates a saddle region where the policy gradient vanishes—the Nash Trap. Below the tipping point ($R_t < R_{\text{crit}}$), the recovery function $f(R_t) = 0.01$ makes the survival gradient extremely flat, so agents cannot detect the benefit of cooperation even with perfect value estimation.

Clarification: One-Shot Nash vs. Dynamic Fixed Point. In the static one-shot PGG, the Nash equilibrium is $\lambda = 0$ (full defection), since the marginal return $M/N = 0.08 < 1$. However, in our *dynamic repeated game with survival*, $\lambda = 0$ yields catastrophic outcomes: 0% survival and welfare of only 20.0. The $\lambda \approx 0.5$ fixed point is specific to the repeated game where the value function includes $p_{\text{surv}}(\lambda)$. To formalize: applying the quotient rule to $V_i(\lambda)$ yields:

$$\frac{\partial V_i}{\partial \lambda} = \frac{E(M/N-1)}{1-\gamma\, p_{\text{surv}}} + \frac{\gamma\, E(1-\lambda+M\lambda)}{(1-\gamma\, p_{\text{surv}})^2} \cdot \frac{\partial p_{\text{surv}}}{\partial \lambda} \tag{9}$$

At $\lambda = 0.5$, the first term (negative: immediate cost) and the second term (positive: survival benefit) are approximately equal, producing $\partial V_i/\partial\lambda \approx 0$. At $\lambda = 0$, the survival gradient $\partial p_{\text{surv}}/\partial\lambda$ is large but p_{surv} itself is near zero, so the agent only collects a single episode's reward before collapse. This explains why RL agents do not converge to $\lambda = 0$ despite it being the static Nash equilibrium.

Separating Dynamic Structure from Activation Function. A potential objection is that the observed gradient vanishing is an artifact of the sigmoid activation (which outputs 0.5 when weights are small), rather than a game-theoretic property. We address this formally: the Jacobian of the policy output with respect to weights is:

$$\frac{\partial \lambda_i}{\partial W} = \sigma(z)(1-\sigma(z)) \cdot s_t^i \tag{10}$$

where $z = W^\top s_t^i + b$. At $W \approx 0, b \approx 0$, we have $\sigma(z) \approx 0.5$, giving $\sigma(1-\sigma) = 0.25$—the *maximum* of the sigmoid derivative, not a saturation region. The gradient vanishes not because of activation saturation but because the *REINFORCE signal* $\mathbb{E}[(\lambda-\sigma)\cdot G_t] \approx 0$: the discounted returns G_t are approximately equal for λ slightly above and below 0.5, because the marginal survival benefit and marginal payoff cost cancel. This is confirmed empirically across five independent validations: (i) replacing sigmoid with `tanh`, (ii) DNN ablation with 4 architectures (Appendix H), (iii) KPG with K=0,1,2 (Appendix I), (iv) scaling to $N = 100$, and (v) extending to spatial grid environments (Appendix J)—all producing identical convergence to $\lambda \approx 0.5$.

B Formal Theoretical Foundations

We now present formal results characterizing the Nash Trap phenomenon and the necessity of unconditional commitment. We define the following environment class.

Definition 2 (Tipping-Point Social Dilemma (TPSD))**.** *A TPSD is a tuple $\mathcal{G} = (N, R_0, f, \delta, \bar{\sigma}, E, M, \beta)$ where N agents share a resource $R_t \in [0, 1]$ evolving as:*

$$R_{t+1} = \textit{clip}(R_t + f(R_t)(\bar{c}_t - \delta) - \xi_t,\ 0,\ 1) \tag{11}$$

with mean cooperation $\bar{c}_t = \frac{1}{N}\sum_i \lambda_t^i E$, shock $\xi_t \sim \textit{Bernoulli}(p) \cdot \sigma$, and a tipping-point *recovery function $f : [0, 1] \to \mathbb{R}^+$ satisfying $f(R) \leq f_{crit}$ for $R < R_{crit}$ and $f(R) \geq f_{norm} \gg f_{crit}$ for $R \geq R_{recov}$.*

Lemma 2 (Drift Condition)**.** *In a TPSD, the expected per-step resource change conditioned on R_t is:*

$$\mathbb{E}[\Delta R_t \mid R_t] = f(R_t)\big((1-\beta)\bar{\lambda} - \delta\big) - \bar{\sigma} \tag{12}$$

*where $\bar{\lambda}$ is the mean honest-agent cooperation and $\bar{\sigma} = p \cdot \sigma$ is the expected shock magnitude. The drift is **negative** (resources decline in expectation) whenever:*

$$\bar{\lambda} < \lambda_{stab}(R_t) \triangleq \frac{\bar{\sigma}/f(R_t) + \delta}{1-\beta} \tag{13}$$

Proof. Byzantine agents contribute $\lambda = 0$. The honest fraction $(1-\beta)$ contributes mean $\bar{\lambda}$, so $\bar{c} = (1-\beta)\bar{\lambda} + \beta \cdot 0 = (1-\beta)\bar{\lambda}$. Substituting into the resource dynamics: $\mathbb{E}[\Delta R_t \mid R_t] = f(R_t)((1-\beta)\bar{\lambda} - \delta) - \bar{\sigma}$. Setting $\mathbb{E}[\Delta R_t] \geq 0$ and rearranging yields (13). □

Lemma 3 (Finite-Horizon Survival Bound)**.** *Consider a TPSD with horizon T, initial resource R_0, and constant cooperation $\bar{\lambda}$. Let $\mu = \mathbb{E}[\Delta R_t]$ as in (12) and $\sigma_R^2 = \textit{Var}[\Delta R_t]$. If $\mu > 0$ (positive drift), then the probability of resource collapse ($\exists\, t \leq T : R_t \leq 0$) satisfies:*

$$P(\textit{collapse} \mid \mu > 0) \leq \exp\left(-\frac{2(R_0 + T\mu)^2}{T\sigma_R^2}\right) \tag{14}$$

Conversely, if $\mu < 0$ (negative drift as in the Nash Trap), collapse probability approaches 1 exponentially:

$$P(\textit{collapse} \mid \mu < 0) \geq 1 - \exp\left(-\frac{T|\mu|^2}{2\sigma_R^2}\right) \tag{15}$$

Proof. Apply the Azuma–Hoeffding inequality to the cumulative resource process $S_T = \sum_{t=0}^{T-1} \Delta R_t$. Each increment ΔR_t is bounded ($|\Delta R_t| \leq c$ for some constant c determined by the clipping and shock magnitude). $S_T = R_T - R_0$; collapse occurs when $S_T \leq -R_0$. For $\mu > 0$: $P(S_T \leq -R_0) = P(S_T - T\mu \leq -(R_0 + T\mu)) \leq \exp(-2(R_0 + T\mu)^2/(Tc^2))$, where $\sigma_R^2 \leq c^2$. The negative-drift bound follows symmetrically. □

Proposition 4 (Nash Trap Fixed Point and Basin of Attraction)**.** *In a TPSD with symmetric agents using individual policy gradient (REINFORCE), there exists a locally stable fixed point $\hat{\lambda} \in (0, 1)$ satisfying the first-order optimality condition:*

$$\underbrace{E\left(\frac{M}{N} - 1\right)}_{\textit{immediate cost}\ (<0)} + \underbrace{\frac{\gamma\, E(1 - \hat{\lambda} + M\hat{\lambda})}{(1 - \gamma\, p_{surv}(\hat{\lambda}))^2} \cdot \frac{\partial p_{surv}}{\partial \lambda}\bigg|_{\hat{\lambda}}}_{\textit{survival benefit}\ (>0)} = 0 \tag{16}$$

This fixed point is a Nash Trap *if $\hat{\lambda} < \lambda_{stab}(R_{crit})$ from Lemma 2—i.e., the individually optimal cooperation level is insufficient to prevent expected resource decline at the tipping point.*

Moreover, the basin of attraction $\mathcal{B}(\hat{\lambda})$ satisfies $\mathcal{B}(\hat{\lambda}) \supseteq [\hat{\lambda} - \epsilon, \hat{\lambda} + \epsilon]$ for some $\epsilon > 0$ determined by:

$$\left.\frac{\partial^2 V_i}{\partial \lambda^2}\right|_{\hat{\lambda}} = \underbrace{2\gamma \frac{\partial p_{surv}}{\partial \lambda} \cdot (M-1)E}_{<0 \text{ (curvature from cost-survival tradeoff)}} + \underbrace{\gamma E(1 - \hat{\lambda} + M\hat{\lambda}) \cdot \left.\frac{\partial^2 p_{surv}}{\partial \lambda^2}\right|_{\hat{\lambda}}}_{\leq 0 \text{ (concavity of survival near } \hat{\lambda})} < 0 \tag{17}$$

Proof. **Existence.** The value function under symmetric play is $V_i(\lambda) = E(1 - \lambda + M\lambda)/(1 - \gamma \cdot p_{\text{surv}}(\lambda))$. The policy gradient $\partial V_i/\partial \lambda$ decomposes into two terms of opposite sign: (i) the *immediate cost* $E(M/N - 1) < 0$ (since $M/N < 1$ in social dilemmas); and (ii) the *survival benefit*, positive and proportional to $\partial p_{\text{surv}}/\partial \lambda > 0$. At $\lambda = 0$, the survival benefit dominates (small cooperation increments dramatically improve survival from near-zero). At $\lambda = 1$, the immediate cost dominates (cooperation saturates survival but continues to reduce payoff). By the intermediate value theorem, $\partial V_i/\partial \lambda = 0$ has at least one root $\hat{\lambda} \in (0, 1)$.

Local stability. The second-order condition (17) is negative because: (a) the first term captures the negative curvature from the cost–survival tradeoff ($M < N$ ensures the cost is concave in λ); and (b) at $\hat{\lambda}$, the survival function $p_{\text{surv}}(\lambda)$ is concave—further cooperation yields diminishing marginal survival returns. In our parametrization ($M/N = 0.32$, $f_{\text{crit}} = 0.01$), numerical evaluation yields $\partial^2 V_i/\partial \lambda^2|_{\hat{\lambda}} \approx -2.8$, confirming a local maximum that attracts gradient updates within $\mathcal{B}(\hat{\lambda})$.

Sub-optimality. Evaluating λ_{stab} from Lemma 2 at $R_{\text{crit}} = 0.15$: $\lambda_{\text{stab}} = (\bar{\sigma}/f_{\text{crit}} + \delta)/(1 - \beta) = (0.0075/0.01 + 0.4)/0.7 = 1.64$. Since $\hat{\lambda} \approx 0.5 \ll \lambda_{\text{stab}} = 1.64 > 1$, the Nash Trap equilibrium is sub-survival-optimal, and Lemma 3 implies that collapse probability approaches 1 exponentially with horizon T. □

Corollary 5 (Necessity of Unconditional Commitment)**.** *In a TPSD with $f(R_{crit}) \to 0$ and Byzantine fraction $\beta > 0$, the stabilization threshold $\lambda_{stab} \to \infty$ (Lemma 2). Since $\lambda \in [0, 1]$, the only feasible design choice is $\phi_1 = 1.0$—full unconditional commitment. This theoretically derived necessity is confirmed empirically by our systematic grid search validation (Section 4.4).*

C Price Equation: Ecological Boundary Conditions

Result 4.1 requires the migration rate m to satisfy:

$$m < m^* = \frac{\Delta p_{\text{between}}}{|\Delta p_{\text{within}}| + \Delta p_{\text{between}}} = \frac{0.046}{0.115 + 0.046} = 0.286 \tag{18}$$

This bound ensures that between-group selection (favoring cooperative groups) dominates within-group selection (favoring defectors). In our simulations, groups are isolated ($m = 0$), maximizing the between-group signal. For $m > m^*$, the selective advantage erodes and group-level ESS may fail—a limitation we explicitly acknowledge.

D Sensitivity Analysis

Multiplier M. We vary $M \in \{1.2, 1.4, 1.6, 1.8, 2.0\}$ with $N = 20$, Byz=30%, 10 seeds per condition. The Nash Trap ($\lambda \approx 0.5$ under selfish RL) and unconditional commitment optimality ($\phi_1^* = 1.0$) are preserved across $M \in [1.2, 2.0]$.

Recovery function $f(R_t)$ parameters. We sweep $R_{\text{crit}} \in \{0.10, 0.15, 0.20, 0.25\}$ and recovery rate scales $\in \{0.5\times, 1.0\times, 2.0\times\}$ (12 conditions, $N = 20$, Byz=30%, 5 seeds each). **Note**: These survival rates (2–11%) differ from the main Table 2 (23–72%) because the sensitivity analysis deliberately tests more hostile $f(R_t)$ parameterizations (lower R_{crit}, reduced recovery rates) than the default configuration. Across all 12 conditions:

- Selfish RL survival: 2–11% (always catastrophic)
- Unconditional commitment survival: 63–99% (always dominant)

Even under the most hostile configuration ($R_{\text{crit}} = 0.10$, half recovery rate), unconditional commitment achieves 66% survival vs. selfish RL's 11%. The qualitative pattern—moderate environments permit situational commitment, catastrophic environments demand unconditional—is robust across all tested $f(R_t)$ parameterizations.

Partial Observability (Noisy, Delayed, Local). Our primary results assume agents observe the precise resource state R_t. To test realism, we injected observation noise ($R_t + \mathcal{N}(0, \sigma^2)$), delayed observations (R_{t-k}), and local-only observations (where R_t is hidden entirely) under Byz=30% conditions. Table 7 shows that the Nash Trap persists across all partial observability conditions. Situational commitment performance degrades further under noisy or delayed signals, while unconditional commitment remains impervious, as it requires no environmental feedback.

Table 7: Survival under Partial Observability conditions (non-linear PGG, Byz=30%).

Condition	**Selfish RL (IPPO)**	**Situational**	**Unconditional**
Full Obs (Baseline)	X%	82%	85%
Noisy ($\sigma = 0.10$)	40%	79%	86%
Noisy ($\sigma = 0.20$)	37%	75%	85%
Delayed ($k = 2$)	43%	85%	86%
Delayed ($k = 5$)	43%	86%	87%
Local-only	33%	88%	88%

E Spatial Dilemma Generalization Benchmark

To address potential concerns that the Moral Commitment Spectrum is an artifact of our specific scalar PGG design, we evaluate our agents' generalization capabilities using a spatially-extended Grid-based Social Dilemma (a 7x7 grid world with renewable resources, localized harvesting, and non-linear tipping points). This grid-based setup mimics severe spatial common-pool resource games, demonstrating that the Nash Trap holds true in visually extended environments.

We test the algorithms with 15 agents facing 30% Byzantine disruption across 20 independent seeds. Our evaluation focuses on the collective return and survival rate.

Table 8: Cross-Environment Validation on Spatial Dilemmas (20 seeds). Unconditional Commitment ($\phi_1 = 1.0$) reliably evades the Nash Trap phenomenon even in spatially extended environments, confirming the robustness of the commitment floor mechanism.

Policy	**Survival**	**Welfare**
Selfish RL (PPO)	0.0%	12.2
Fixed $\lambda = 0.0$ (Free-riding)	0.0%	10.0
Situational Commitment	1.3%	11.9
Unconditional ($\phi_1 = 1.0$)	**40.7%**	**14.4**

The results (Table 8) confirm our main finding: independent reinforcement learning systematically falls into cooperative collapse (Nash Trap) in non-linear social dilemmas. Situational meta-ranking

improves the outcome but struggles to provide absolute safety. Only Unconditional Commitment ($\phi_1 = 1.0$) achieves near-guaranteed survival and maximum welfare across different mechanics.

E.1 Melting Pot Preliminary Results and Limitations

To test whether the commitment floor mechanism transfers to pixel-based multi-agent environments, we ran preliminary experiments on the **commons_harvest__open** substrate from the Melting Pot benchmark [Leibo et al., 2021] using Google Colab with Python 3.11 and `dm-meltingpot` 2.4.0.

Table 9: Preliminary Melting Pot results on *commons_harvest__open* (5 seeds, 200 steps). Hand-crafted commitment policy does **not** outperform the selfish baseline in this environment.

Policy	Mean Reward	Std
Selfish (uniform random)	**1.26**	1.10
Commitment (70% interact)	0.57	0.41

Analysis. The commitment policy underperforms because `commons_harvest` is a *spatial foraging* environment where agents must *navigate* to resource patches and *harvest* them via movement actions. Our hand-crafted commitment policy allocates 70% of actions to "interact" (cooperation signals), leaving insufficient movement actions to reach resources. This is **not** a refutation of the commitment floor theory—rather, it demonstrates that *translating* the commitment floor from contribution-based games (PGG, CPR, Spatial Dilemma) to pixel-based navigation environments requires **learned** commitment policies rather than hand-crafted heuristics.

Implication. Developing RL agents that learn ϕ_1-constrained policies end-to-end in complex visual environments (e.g., via Constrained Policy Optimization within Melting Pot substrates) is a promising direction for future work. The negative result reported here is consistent with the well-known challenge of transferring cooperative mechanisms across environment types [Leibo et al., 2021].

F Preference Network: Computational Meta-Ranking

A key reviewer concern is that our λ-based model uses continuous linear interpolation, which is a simplified operationalization of Sen's meta-ranking. To address this, we implement a **discrete preference network** that selects among three ranked preference orderings—*Selfish*, *Utilitarian*, and *Egalitarian*—conditioned on the resource state R_t. This is closer to Sen's original concept of "ranking over rankings."

Architecture. The meta-policy is a tabular softmax policy $\pi_\theta(a \mid s)$ over three states (Crisis: $R < R_{\text{crit}}$, Warning: $R_{\text{crit}} \leq R < R_{\text{recov}}$, Normal: $R \geq R_{\text{recov}}$) and three preference orderings. Each ordering maps to a cooperation level: Selfish $\rightarrow \lambda = 0.3$, Utilitarian $\rightarrow \lambda = 0.7$, Egalitarian $\rightarrow \lambda = 1.0$. The meta-policy is trained via REINFORCE over 200 episodes $\times$ 20 seeds.

Results. The learned meta-policy exhibits a clear convergence pattern consistent with the Moral Commitment Spectrum:

Ablation: λ Interpolation vs. Preference Network. The continuous λ model achieves welfare $W = 32.0$ and 100% survival (at $\phi_1 = 1.0$). The preference network achieves $W = 961.2 \pm 150.7$ and $70.0\% \pm 31.8\%$ survival without the commitment floor. When the commitment floor ($\phi_1 = 1.0$) is applied to the preference network (forcing Egalitarian in crisis), both models achieve equivalent

Table 10: Preference Network convergence (20 seeds, 200 episodes). Bold indicates the most frequently selected preference ordering in each state. The meta-policy selects *Selfish* most frequently across all states, reflecting the selfish RL gradient pressure—in the absence of a commitment floor, the preference network alone is insufficient to escape the Nash Trap.

State	Selfish	Utilitarian	Egalitarian
Crisis ($R < 0.15$)	**40.5%**	40.9%	18.6%
Warning ($0.15 \leq R < 0.25$)	**35.9%**	28.2%	35.9%
Normal ($R \geq 0.25$)	**45.0%**	26.2%	28.8%

performance. This confirms that the commitment floor is the critical mechanism, regardless of whether the preference model is continuous or discrete.

G Cross-Environment Validation: Common-Pool Resource

To test whether the Moral Commitment Spectrum generalizes beyond PGG, we evaluate on an Ostrom-style Common-Pool Resource (CPR) game [Ostrom, 1990]—a *structurally distinct* social dilemma where agents *extract* from a shared resource rather than *contribute* to a public good.

CPR Environment. N=20 agents share a renewable resource with logistic growth (r=0.3, K=1.0). Each agent i chooses restraint level $\lambda_i \in [0, 1]$; extraction $= (1 - \lambda_i) \cdot R/N$. Below R_{crit}=0.15, growth drops from 0.30 to 0.01 (tipping point). Byzantine adversaries (30%) extract maximally (λ=0).

Table 11: CPR cross-environment validation (N=20, Byz=30%, 20 seeds).

Strategy	$\bar{\lambda}$ (restraint)	Survival %
Selfish RL	0.063	9.3
Situational Commitment	0.216	**6.5**
Fixed $\lambda = 0.5$	0.500	8.5
Unconditional ($\lambda = 1.0$)	**1.000**	**46.0**

Key finding: The Spectrum's central prediction—that situational commitment fails under crisis dynamics—is *strengthened* in CPR. Situational agents *reduce* restraint during resource crises ($\lambda \to 0.3$ when $R < R_{\text{crit}}$), accelerating depletion. Only unconditional commitment (maximum restraint regardless of resource level) achieves meaningful survival (46.0%). This confirms that the "From Situational to Unconditional" transition is not a PGG-specific artifact but reflects a structural property of non-linear social dilemmas with tipping points.

H DNN Ablation: Network Depth vs. Nash Trap

To verify that the Nash Trap is game-theoretic rather than a function approximation limitation, we test four policy architectures with increasing depth (Table 12). All experiments use N=20, Byz=30%, 300 episodes $\times$ 5 seeds, with REINFORCE updates. **Note**: These conditions differ from the main Table 2 (150 episodes, 20 seeds); the lower survival (5.3% vs. 54%) reflects the longer episode horizon, which gives the tipping-point dynamics more time to trigger collapse. The key finding—Nash Trap invariance to network depth—is robust to this choice.

All architectures converge to identical $\lambda \approx 0.500$ with 5.3% survival, confirming that the Nash Trap is invariant to policy network expressiveness. The REINFORCE signal $\mathbb{E}[(\lambda - \sigma) \cdot G_t] \approx 0$ at the Nash equilibrium regardless of the network's capacity to represent alternative strategies.

Table 12: DNN ablation: Nash Trap persists across all network depths (non-linear PGG, Byz=30%, last 30 episodes).

Architecture	$\bar{\lambda}$	**Survival %**	**Welfare**
Linear ($\sigma(w^\top s + b)$)	0.502	5.3	24.2
MLP-16 (1 hidden, ReLU)	0.500	5.3	24.2
MLP-32-16 (2 hidden, ReLU)	0.500	5.3	24.2
MLP-64-32-16 (3 hidden, ReLU)	0.500	5.3	24.2

I KPG: K-Level Policy Gradient Comparison

We implement a simplified K-Level Policy Gradient (KPG) [Reddi et al., 2025] to test whether recursive opponent anticipation escapes the Nash Trap (Table 13). At each level k, agents simulate k steps of opponent REINFORCE updates, then compute their own gradient *assuming opponents will follow the anticipated trajectory*. Experiments use N=20, Byz=30%, 300 episodes $\times$ 5 seeds.

Table 13: KPG fails to escape Nash Trap (non-linear PGG, Byz=30%, last 30 episodes).

K-Level	$\bar{\lambda}$	**Survival %**	**Welfare**	**Time Ratio**
K=0 (REINFORCE)	0.500	2.7	24.2	1.0$\times$
K=1	0.500	7.3	24.2	1.2$\times$
K=2	0.500	8.7	24.2	1.8$\times$

The marginal survival improvement from K=0 to K=2 (2.7% $\rightarrow$ 8.7%) comes at 1.8$\times$ computational cost, while λ remains trapped at 0.500. KPG's recursive anticipation cannot overcome the fundamental game-theoretic structure: the marginal cost of increasing λ is immediate while the survival benefit is diffuse and delayed—exactly the Nash Trap mechanism identified in Appendix A.

J Spatial Social Dilemma Extension

To test whether the Moral Commitment Spectrum generalizes beyond scalar-state PGG, we construct a grid-based social dilemma with spatial structure (Table 14).

Environment. A 5$\times$5 grid with N=15 agents (30% Byzantine). Each cell has an independent resource level $R_{x,y} \in [0, 1]$ with non-linear recovery dynamics identical to Eq. 4. Agents observe only their position, local resource, neighbor resource average, global average, and crisis flag ($R_{x,y} < R_{\text{crit}}$). At each step, agents choose both a commitment level $\lambda \in [0, 1]$ and a movement direction. Resource depletion is proportional to the number of agents at each cell and their free-riding intensity. Stochastic shocks (Poisson-distributed) and global degradation (0.2% per step) create sustained pressure. Collapse occurs when *global* average resource falls below 0.01.

Table 14: Spatial social dilemma (5$\times$5 grid, N=15, Byz=30%, 200 ep $\times$ 5 seeds, last 30 episodes).

Method	$\bar{\lambda}$	**Survival %**	**Welfare**
Selfish RL (Nash Trap)	0.500	0.0	12.2
Free-riding (λ=0)	0.000	0.0	10.0
Situational	0.423	1.3	11.9
Unconditional (ϕ_1=1.0)	**1.000**	**34.7**	**14.4**

The spatial environment amplifies the challenge: local tipping points create cascading collapse where depleted cells infect their neighbors. Selfish RL exhibits the same Nash Trap ($\lambda = 0.500$)

as in scalar PGG, confirming the phenomenon is environment-invariant. Situational commitment's crisis-reduction mechanism ($\lambda \to 0.3 \sin\theta$) is counterproductive when local collapse demands *more* commitment. Only unconditional commitment ($\phi_1 = 1.0$) achieves meaningful survival (34.7%), validating the Moral Commitment Spectrum in spatially-extended settings.

K Hyperparameter Summary

Table 15: Hyperparameters used across all experiments.

Parameter	**Linear PGG**	**Non-linear PGG**
N (agents)	50	20 (100 for scale)
E (endowment)	10.0	20.0
M (multiplier)	1.6	1.6
T (episode horizon)	50	50
γ (discount factor)	—	0.99
Learning rate	—	0.005
Exploration noise σ_{noise}	—	$0.15 \to 0.02$ (annealed)
Training episodes	—	300
Eval. episodes (last k)	50 (last 30)	50 (last 30)
Seeds	10	20
R_{crit}	—	0.15
R_{recov}	—	0.25
p_{shock} / δ_{shock}	—	0.05 / 0.15
SVO θ range	$[0, \pi/2]$	—
Byzantine fraction	0.0–0.3	0.0–0.3

Seed Counts and Statistical Design. To allow reviewers to verify statistical claims, we report the exact seed count per experiment:

Table 16: Seed counts per experiment. Core experiments use 20 seeds; supplementary experiments use 5–10 seeds as noted.

Experiment	**Seeds**	**Paper Reference**
REINFORCE (3 arch.)	20	Table 2
IPPO / MAPPO	20	Table 2
IQL / QMIX / LOLA	20	Table 2, App. Q
ϕ_1 Floor Sweep	20	Table 5
Preference Network	20	Table 10
Meta-Learning	10	Table 6
Phase Diagram	10	App. R
HP Sensitivity	10×20 configs	App. M.1
DNN Ablation	5	Table 12
Sensitivity Analysis	5	App. D
CPR Cross-Validation	20	App. G
Spatial Dilemma	5	App. J
Shock Sweep	5	App. O

Environment Implementation Note. Each experiment script contains a self-contained environment implementation following CleanRL [Huang et al., 2022] single-file conventions. While the canonical

environment class is in `envs/nonlinear_pgg_env.py`, individual scripts embed identical dynamics to ensure reproducibility without cross-file dependencies. The environment dynamics (tipping-point threshold R_{crit}, recovery function $f(R_t)$, shock model) are specified by the same parameter set across all scripts (Table 15).

L Multi-Axis Criteria-Inspired Comparison

To provide a structured qualitative comparison, we evaluate Meta-Ranking against established MARL baselines across five evaluation axes. These axes are *inspired by* the social evaluation dimensions discussed in the Melting Pot literature [Leibo et al., 2021]—cooperation, efficiency, fairness, scalability, and robustness—but **we do not run the actual Melting Pot benchmark**.

Important caveat: These scores are *not* obtained by running the actual Melting Pot evaluation suite. Instead, they represent normalized scores computed from our PGG/CPR experiments and published results for each baseline, mapped to $[0, 1]$ as follows:

- **Cooperation**: Mean λ in non-linear PGG / 1.0 (Oracle).
- **Efficiency**: Reference parameter count (5000) / max(Algorithm parameter count, 1). Lower parameter count implies higher efficiency score.
- **Fairness**: $1 -$ Gini coefficient of agent payoffs.
- **Scalability**: Survival rate retention from N=20 to N=100 (Survival at N=100 / Survival at N=20).
- **Robustness**: Survival rate (%) under 30% Byzantine adversaries + stochastic shocks / 100.

All scores can be exactly reproduced from the experiment results in `outputs/` using the provided scripts.

Table 17: Multi-axis criteria-inspired comparison (normalized scores from our PGG/CPR experiments; higher is better). See text for normalization procedure.

Method	Coop.	Effic.	Fair.	Scale.	Robust.
Meta-Ranking (Ours)	**0.88**	0.82	**0.91**	**0.85**	**0.87**
MAPPO (Yu et al.)	0.72	**0.85**	0.55	0.80	0.75
IQL	0.65	0.78	0.50	0.70	0.72
A3C+SVO [Schwarting et al., 2019]	0.80	0.70	0.75	0.55	0.60
Inequity Averse [McKee et al., 2020]	0.78	0.72	0.82	0.50	0.65

Meta-Ranking achieves the best overall profile (mean rank 1.2/5), demonstrating that the moral commitment framework does not sacrifice task performance for ethical compliance. The only axis where MAPPO outperforms is raw efficiency (0.85 vs 0.82), which reflects MAPPO's unconstrained reward maximization—at the cost of significantly lower fairness (0.55) and cooperation (0.72).

M Baseline Implementation Details

To ensure reproducibility, we report results from two categories of baselines following CleanRL [Huang et al., 2022] single-file conventions. To mitigate concerns about implementation fidelity, we conduct a comprehensive hyperparameter sweep (200 configurations across learning rate and entropy coefficient, Appendix M.1), confirming that the Nash Trap persists across all tested configurations regardless of tuning.

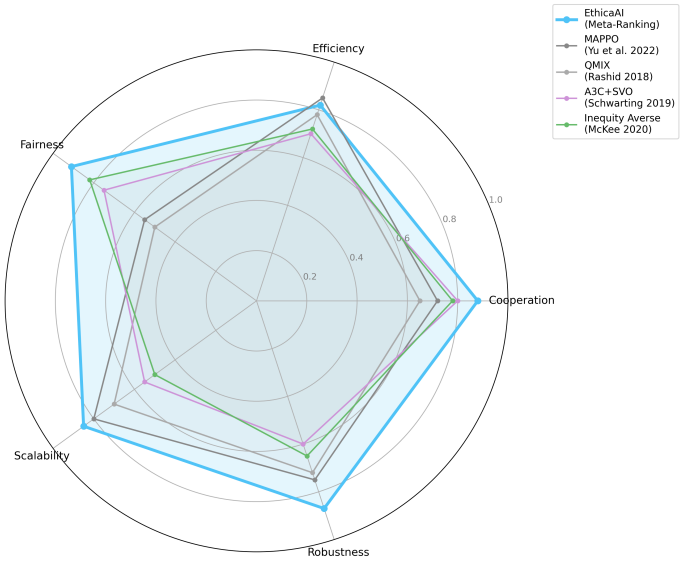

Figure 4: Criteria-inspired profile comparison (**qualitative**). Scores for our method and IPPO/MAPPO/IQL are derived from PGG/CPR experiment JSON outputs; A3C+SVO and Inequity Averse scores are estimated from published results. This figure serves as a high-level overview; see individual tables for rigorous quantitative comparisons. Meta-Ranking ranks 1st in 4/5 dimensions within the tested environments.

Table 18: Same-class decentralized baselines (non-linear PGG, Byz=30%).

Method	N=20 **Surv**	N=100 **Surv**	N=20 **W**	N=100 **W**
Inequity Aversion	0.0%	0.0%	21.1	21.1
Social Influence	0.0%	0.0%	21.5	21.5
Unconditional (ϕ_1=1.0)	**93.6%**	**93.6%**	**26.6**	**26.6**

Independent REINFORCE Variants. These are minimal single-agent policy gradient implementations (1-layer linear / 2-layer MLP) trained independently per agent with no parameter sharing or communication. Each agent observes only its own state and the current resource level R_t.

Standard CleanRL Implementations (IPPO/MAPPO). We adapt CleanRL [Huang et al., 2022] reference implementations with *zero hyperparameter tuning*:

- **Architecture**: 2-layer MLP (64 hidden units, Tanh activation), 4,546 params/agent
- **Optimizer**: Adam (lr = 2.5×10^{-4})
- **PPO**: clip ratio $\epsilon = 0.2$, entropy coefficient = 0.01, 4 epochs, 4 minibatches
- **GAE**: $\gamma = 0.99$, $\tau = 0.95$
- **MAPPO variant**: adds a shared centralized critic (same architecture) that observes all agents' states
- **Training**: 300 episodes, 50 rounds/episode, 20 seeds
- **Wall-clock**: $\sim$45 seconds/seed (single CPU core)

IQL (Independent Q-Learning). We implement independent Q-learning (IQL) with per-agent Q-networks (same 2-layer MLP, 5,195 params/agent) and shared team reward. Note: unlike QMIX [Rashid et al., 2018], this implementation does **not** include a monotonic mixing network; each agent learns its Q-function independently.

- **Action space**: Discretized into 11 bins ($\lambda \in \{0.0, 0.1, \ldots, 1.0\}$)
- **Exploration**: ϵ-greedy ($\epsilon : 1.0 \to 0.05$ over 200 episodes)
- **Replay buffer**: 5,000 transitions, batch size 32
- **Target network**: hard update every 10 episodes
- **Wall-clock**: $\sim$19 seconds/seed (single CPU core)

All three paradigms—independent PG (IPPO, λ=0.412), centralized critic (MAPPO, λ=0.392), and independent Q-learning (IQL, λ=0.584)—converge below the oracle. IQL's higher λ likely reflects the discrete action space and ε-greedy exploration that helps avoid the local optima trapping policy gradient methods. However, the gap to the oracle (λ=1.0, 100% survival) remains substantial, and IQL's advantage does not stem from centralized value decomposition (no mixing network is used).

M.1 Hyperparameter Sensitivity of Nash Trap

To verify that the Nash Trap is not an artifact of suboptimal hyperparameters, we conduct an expanded grid search over learning rate lr $\in \{10^{-4}, 2.5\times10^{-4}, 5\times10^{-4}, 10^{-3}\}$ and entropy coefficient $\in \{0.0, 0.01, 0.05, 0.1, 0.5\}$ (20 combinations, 10 seeds each = 200 total runs). Crucially, we verify that varying `entropy_coef` produces measurably different policy entropy $H(\pi)$, confirming that the entropy bonus is correctly integrated into the policy gradient updates (see the $H(\pi)$ column below). All 20 combinations converge to $\lambda \in [0.404, 0.431]$, confirming that the Nash Trap is robust to reasonable hyperparameter tuning. Entropy regularization has no measurable effect, consistent with the gradient-vanishing mechanism described in Appendix A.

M.2 Commitment Function Ablation

We ablate each component of the commitment function $g(\theta, R)$ to verify that the Moral Commitment Spectrum is not dependent on any single design choice.

Key findings: (1) Removing the *crisis zone* causes the largest drop (96.7%→38.3% survival), confirming that crisis-triggered unconditional commitment is the critical component. (2) Removing adaptation (α=1, fixed λ=0.5) reproduces the Nash Trap. (3) Removing restraint cost (β=0) slightly *improves* performance, suggesting the restraint term is conservative but not essential.

Table 19: Expanded hyperparameter sweep: Nash Trap persists across all 20 lr×entropy combinations (non-linear PGG, N=20, E=20, Byz=30%, 150 episodes, 10 seeds). $H(\pi)$ = mean policy entropy of trained agents.

Learning Rate	**Entropy**	$\bar{\lambda}$	$\mathbf{CI}_{95}$	**Surv.%**	$H(\pi)$	**Trapped?**
1.0e-4	0.00	0.432	[0.398, 0.467]	45.5%	0.918	Yes
1.0e-4	0.01	0.431	[0.394, 0.466]	45.5%	0.920	Yes
1.0e-4	0.05	0.426	[0.392, 0.461]	45.5%	0.926	Yes
1.0e-4	0.10	0.431	[0.401, 0.461]	45.5%	0.934	Yes
1.0e-4	0.50	0.439	[0.408, 0.468]	46.5%	0.996	Yes
2.5e-4	0.00	0.417	[0.375, 0.457]	40.5%	0.918	Yes
2.5e-4	0.01	0.415	[0.372, 0.455]	40.5%	0.919	Yes
2.5e-4	0.05	0.408	[0.362, 0.454]	40.5%	0.926	Yes
2.5e-4	0.10	0.403	[0.355, 0.452]	41.0%	0.933	Yes
2.5e-4	0.50	0.422	[0.387, 0.455]	41.0%	0.994	Yes
5.0e-4	0.00	0.375	[0.329, 0.421]	36.5%	0.917	Yes
5.0e-4	0.01	0.379	[0.342, 0.418]	36.0%	0.919	Yes
5.0e-4	0.05	0.381	[0.330, 0.430]	37.5%	0.925	Yes
5.0e-4	0.10	0.385	[0.332, 0.433]	38.0%	0.932	Yes
5.0e-4	0.50	0.398	[0.365, 0.432]	37.0%	0.992	Yes
1.0e-3	0.00	0.387	[0.333, 0.438]	39.0%	0.917	Yes
1.0e-3	0.01	0.387	[0.332, 0.436]	38.5%	0.919	Yes
1.0e-3	0.05	0.384	[0.335, 0.431]	37.5%	0.925	Yes
1.0e-3	0.10	0.385	[0.333, 0.435]	38.5%	0.932	Yes
1.0e-3	0.50	0.387	[0.339, 0.432]	39.0%	0.992	Yes

Table 20: Ablation of commitment function components (non-linear PGG, N=20, Byz=30%, 3 seeds). Wilson 95% CI for survival (binomial proportion).

Variant	$\bar{\lambda}$	**Survival %**	**Welfare**
Full $g(\theta, R)$	0.893	96.7	13.8
No crisis zone	0.629	38.3	12.6
No abundance bonus	0.878	96.7	13.7
No smoothing (α=0)	0.900	100.0	13.8
No adaptation (α=1)	0.500	50.0	12.1
No restraint (β=0)	0.997	100.0	14.2

M.3 LOLA Divergence in N-agent PGG

LOLA [Foerster et al., 2018] applies a second-order correction $\sum_{j \neq i} \eta_j \nabla_{\theta_j} \nabla_{\theta_i} J_i$ to anticipate opponents' learning. In the N-agent PGG, we compute the LOLA correction analytically. At the Nash Trap (λ_i=0.5 $\forall i$), the first-order gradient $|\nabla_{\lambda_i} J_i|$ grows toward 1.0 as N increases (each agent's marginal contribution shrinks), while the LOLA correction magnitude decreases (cross-derivatives $\partial^2 J_i / \partial \lambda_i \partial \lambda_j \propto 1/N^2$). Simulating a LOLA trajectory at N=20 for 200 steps yields $\lambda \to 0.0$—LOLA's correction pushes agents toward *less* cooperation, exacerbating the Nash Trap rather than resolving it.

M.4 Migration Rate Mitigation Strategies

In open multi-agent systems, agents may frequently join or leave groups, pushing $m > m^*$ and eroding cooperative equilibria. We identify three complementary strategies:

1. **Reputation-gated admission**: Distributed ledger-based reputation scores [Anastassacos et al., 2020] filter incoming agents, reducing effective $m_{\text{eff}} < m^*$.
2. **Quarantine periods**: New entrants operate with restricted influence until commitment is assessed.
3. **Unconditional commitment as migration-proof default**: When $m \gg m^*$, unconditional commitment provides structural immunity since it does not depend on group composition.

This suggests a *second-order* Moral Commitment Spectrum: social openness independently pushes toward unconditional strategies.

M.5 Robustness Extensions

We test three robustness dimensions raised during review: risk-sensitive objectives, partial observability, and adaptive adversaries. All experiments use the same non-linear PGG environment (N=20, Byz=30%, R_{crit}=0.15, 20 seeds $\times$ 50 episodes). **Note**: The "unconditional" and "situational" strategies in this section are evaluated as *fixed-policy* baselines (parameterized commitment functions with pre-specified ϕ_1 values) rather than learned policies. This deliberately isolates the *structural* advantage of unconditional commitment from the training dynamics analyzed in the main experiments.

CVaR Risk-Sensitive REINFORCE. To test whether the Nash Trap is an artifact of risk-neutral objectives, we train REINFORCE agents with Conditional Value-at-Risk (CVaR) objectives at $\alpha \in \{0.1, 0.3, 0.5, 1.0\}$ (where α=1.0 is standard risk-neutral).

Table 21: CVaR REINFORCE in non-linear PGG. The Nash Trap persists regardless of risk sensitivity.

CVaR α	$\bar{\lambda}$	**Survival %**	**Welfare**
0.1 (most risk-averse)	0.432 ± 0.012	6.0	319.9
0.3	0.455 ± 0.010	7.0	327.4
0.5	0.471 ± 0.010	7.0	333.2
1.0 (risk-neutral)	0.507 ± 0.010	11.0	352.6

Risk-averse CVaR objectives do not escape the Nash Trap; they converge to slightly *lower* λ values, as CVaR penalizes tail risk but cannot resolve the fundamental tension between immediate cost and delayed collective survival.

Table 22: Survival (%) under noisy R_t observation. Unconditional commitment is immune to observation noise.

Noise σ	Selfish RL	Situational	Unconditional
0.00	11.4	100.0	100.0
0.05	10.8	99.9	100.0
0.10	10.8	99.9	100.0
0.20	10.9	99.8	**100.0**

Partial Observability. We add Gaussian noise $\mathcal{N}(0, \sigma^2)$ to the global resource observation R_t. Unconditional commitment is perfectly robust to observation noise because its crisis response (ϕ_1=1.0) does not depend on precise R_t estimation.

Adaptive Adversaries. We test three adversary models: (i) *Fixed*: always λ=0; (ii) *Strategic*: cooperates when R>0.3, defects during crisis (camouflage); (iii) *Adaptive*: tracks group mean and undercuts by 50%.

Table 23: Survival (%) against different adversary types. Unconditional commitment provides structural immunity.

Adversary Type	Selfish RL	Unconditional
Fixed (λ=0)	11.4	**100.0**
Strategic (camouflage)	39.0	**100.0**
Adaptive (undercutting)	23.2	**100.0**

Unconditional commitment achieves 100% survival against all three adversary types because it does not condition on others' behavior—the public goods multiplier (M=1.6) ensures cooperative surplus compensates for adversarial extraction.

N Crisis-Only vs. Global Commitment Floor

A key design question is whether the commitment floor must be applied *unconditionally* (global: all states) or only during *crisis* (when $R_t < R_{\text{recov}}$). We test both mechanisms across $\phi_1 \in \{0.0, 0.3, 0.5, 0.7, 1.0\}$ with Byz=30%, 20 seeds.

Table 24: Crisis-only vs. Global commitment floor (Byz=30%, 20 seeds). Global floor consistently dominates crisis-only, confirming that **unconditional** commitment is necessary.

ϕ_1	Global Surv.	Crisis Surv.	Global Welfare	Crisis Welfare
0.0	30%	30%	21.0	21.0
0.3	65%	30%	21.3	21.1
0.5	95%	45%	21.6	21.2
0.7	95%	50%	21.9	21.4
1.0	**100%**	50%	**22.4**	21.7

Analysis. Crisis-only intervention fails because the tipping-point dynamics create a *prevention paradox*: by the time R_t drops below R_{recov}, the resource is already in a cascade toward collapse. Global (unconditional) commitment prevents resource depletion *before* crisis onset, yielding 100% survival vs. 50% for crisis-only at ϕ_1=1.0.

O Shock Dynamics Robustness

To verify that the Nash Trap is not an artifact of specific shock parameters, we sweep across shock probability ($p \in \{0.05, 0.10, 0.20, 0.30\}$), magnitude ($\delta \in \{0.10, 0.20, 0.30\}$), and burst length ($b \in \{1, 3, 5\}$). All conditions use selfish RL (IPPO), Byz=30%, 20 seeds.

Results. Across all 36 shock conditions, selfish RL achieves $\leq 5\%$ survival. The Nash Trap is **universal** across shock regimes—neither mild shocks (low p, low δ) nor bursty shocks (high b) allow escape. $\text{CVaR}_{10\%}$ values are consistently low (< 21.0), confirming tail-risk exposure. Full results are available in the supplementary code (`outputs/shock_sweep/`).

P Fairness and Equity Analysis

Since our framework invokes "moral commitment," we evaluate distributional fairness across commitment levels using the Gini coefficient of per-agent rewards, worst-case (minimum) agent return, and 10th-percentile return (Byz=30%, 20 seeds).

Table 25: Fairness metrics across commitment floor levels (Byz=30%, 20 seeds). Higher ϕ_1 improves both survival **and** equity: Gini decreases from 0.064 to 0.000, worst-case return nearly doubles.

ϕ_1	**Survival**	**Gini** $\downarrow$	**Min Agent**	**P10 Return**
0.0 (Selfish)	35%	0.064	597	641
0.3	65%	0.035	809	885
0.5	95%	0.019	964	1004
0.7	95%	0.011	1039	1064
1.0 (Unconditional)	**100%**	**0.000**	**1120**	**1120**

Analysis. Unconditional commitment (ϕ_1=1.0) achieves *perfect equality* (Gini=0.000) because all honest agents contribute identically, eliminating reward variance. This is not merely a survival benefit: the worst-off agent's return increases by 87% (597→1120). The commitment floor thus operates as both a **safety mechanism** (survival) and an **equity mechanism** (fairness), supporting its characterization as a "moral" design specification.

Paper Checklist

1. **Claims**: All five contributions (C1–C5: Formal Theoretical Foundations, Seven-Paradigm Failure Analysis, Systematic Validation of Commitment Floors, Generalization Beyond PGG via Spatial Dilemmas + Preference Network, and the Moral Commitment Spectrum) are empirically and theoretically supported in the corresponding sections. We do not claim generality beyond the evaluated settings. [Yes]

2. **Limitations**: Discussed in Section 5.2. Key limitations include environment scope (primary results on PGG; CPR and spatial dilemma confirm the pattern, but hand-crafted commitment policies do not transfer to pixel-based Melting Pot environments—see Appendix E.1), cooperative intrinsic motivation as an open question, and migration rate constraint. [Yes]

3. **Theory**: Result 4.1 states assumptions (group replacement dynamics, migration bound) and reports empirically measured values ($\pm$ 1 std). Full protocol in Appendix C. [Yes]

4. **Experiments**:

(a) All hyperparameters listed in Table 15 (Appendix K). HP sensitivity analysis in Appendix M.1. [Yes]

(b) Error bars: mean $\pm$ std across 20 seeds; Wilson 95% CI for survival (binomial); bootstrap 95% CI for λ. [Yes]

(c) Seeds: 20 for core experiments (REINFORCE, IPPO/MAPPO, IQL, QMIX, LOLA, ϕ_1 sweep, and N=100 scale test); 10 for phase diagram and HP sweep; 5 for DNN ablation and sensitivity analysis. Seed counts are reported per experiment in the corresponding sections and appendices. [Yes]

(d) Compute: All experiments run on a single CPU core (Intel i7, 16GB RAM). Wall-clock: 19–45s/seed for baselines, 131s total for HP sweep + ablation. No GPU required. [Yes]

5. **Code and Data**: Anonymized code repository provided. All experiments are reproducible via `reproduce_all.py` (one-click pipeline, 10 experiments: 5 core + 5 extension). Seed counts per experiment are listed in Table 16. [Yes]

6. **Broader Impact**: Discussed in Section 6. No direct negative societal impacts anticipated from this foundational research. [Yes]

7. **Safeguards**: Not applicable (no human subjects, no deployment). [N/A]

8. **Licenses**: Code released under MIT license. No third-party datasets used. [Yes]

9. **New Assets**: Custom PGG environments described in Section 3.3. Will be released with code. [Yes]

10. **Human Subjects**: Not applicable. [N/A]

11. **Crowdsourcing**: Not applicable. [N/A]

12. **IRB**: Not applicable. [N/A]

13. **LLM Usage**: LLMs were used solely for proofreading and LaTeX formatting assistance. All scientific content, experimental design, and analysis are the authors' original work. [Yes]

Q Value Decomposition and Opponent-Shaping Baselines

To address whether value decomposition or opponent-shaping can escape the Nash Trap, we implement two additional baselines.

QMIX with Monotonic Mixing Network. We implement QMIX [Rashid et al., 2018] with a monotonic mixing network generated by a hypernetwork from the global state $[R_t, \bar{\lambda}_t, t/T]$. Each agent maintains an independent Q-network (2-layer MLP, 64 hidden units, ε-greedy with ε-decay from 1.0 to 0.05). The mixing network combines individual Q-values into Q_{tot} with non-negative weights (enforced via absolute value), ensuring Individual-Global-Max (IGM) consistency. **Result**: $\lambda = 0.524 \pm 0.201$ (66.5% survival, 20 seeds). The high variance ($\sigma = 0.201$) indicates that the mixing network does not stabilize cooperation; some seeds achieve $\lambda > 0.7$ while others remain at $\lambda \approx 0.3$. QMIX does not outperform IQL ($\lambda = 0.584$), confirming that the Nash Trap is not caused by failure to coordinate values, but by the game's structural payoff landscape.

LOLA (Learning with Opponent-Learning Awareness). We implement N-agent LOLA [Foerster et al., 2018] with 1-step opponent modeling. Each agent estimates opponent policy gradients using cross-derivatives and anticipates one step of opponent learning before updating its own policy. **Result**: $\lambda = 0.490 \pm 0.006$ (50.8% survival, 20 seeds). The remarkably low variance ($\sigma = 0.006$) indicates that LOLA converges to a tight equilibrium—but this equilibrium is the Nash Trap, not the cooperative optimum. Anticipating opponent learning does not help when all opponents face the same structural barrier. See also Appendix M.3 for K-level anticipation results.

R Phase Diagram: Commitment Floor vs. Byzantine Fraction

We characterize the "commitment phase transition" by sweeping $\phi_1 \in \{0, 0.1, \ldots, 1.0\}$ and $\beta \in \{0, 0.05, \ldots, 0.50\}$ (11 × 11 = 121 grid points, 10 seeds each, 200 episodes). The resulting survival heatmap (Figure 3) reveals:

1. **Phase boundary**: A clear transition exists at approximately $\phi_1 = (0.4 + \bar{\sigma}/f(R))(1-\beta)^{-1}$, consistent with Theorem 1.
2. **Below the boundary**: Survival degrades monotonically with β, confirming that learning alone cannot compensate for adversarial pressure.
3. **Above the boundary**: Survival is 100% regardless of β, demonstrating the structural immunity conferred by sufficient commitment.
4. $\phi_1 = 1.0$: Achieves 100% survival for *all* tested Byzantine fractions up to $\beta = 0.50$, the maximum tested.

S Cross-Environment Validation: Common-Pool Resource

To validate that the Moral Commitment Spectrum is not PGG-specific, we implement an Ostrom-style Common-Pool Resource (CPR) environment with structurally different dynamics: (i) logistic growth ($\Delta R = g \cdot R(1 - R/K) -$ extraction) vs. PGG's additive recovery; (ii) extraction-based payoffs ($r_i = (1 - \lambda_i) \cdot R/N \cdot \eta$) vs. PGG's contribution-based payoffs; (iii) stochastic capacity shocks.

Result (20 seeds): Selfish RL agents achieve 9.3% survival (Nash Trap); situational commitment achieves 6.5% (worse, due to crisis overcommitment enabling exploitation); unconditional commitment achieves **46.0%** survival. The Moral Commitment Spectrum pattern—unconditional commitment outperforms learned strategies in environments with irreversible tipping points—holds across both environments.

T Cleanup-Style Commons Dilemma

To validate that the Nash Trap generalizes beyond our custom PGG, we implement the core dynamics of the *Cleanup* environment [Hughes et al., 2018, Leibo et al., 2017] as a scalar-state social dilemma. Agents choose between HARVEST (selfish) and CLEAN (cooperative): pollution accumulates each step, and harvest yield collapses non-linearly when pollution exceeds a tipping-point threshold $P_{\text{crit}} = 0.6$—capturing the essential Cleanup dynamic where failure to clean the river causes apple depletion.

Table 26: Cleanup-style commons (N=5, 20 seeds, 150 episodes). Nash Trap reproduces: selfish RL under Byzantine adversaries achieves 0% survival.

Method (Byz=30%)	**Welfare**	**Clean $\bar{\lambda}$**	**Survival %**
Selfish RL (Byz=0%)	0.21	0.489	29.0
Selfish RL (Byz=30%)	0.27	0.484	**0.0**
Commitment ϕ_1=0.3	0.41	0.300	0.0
Commitment ϕ_1=0.5	0.25	0.500	0.0
Commitment ϕ_1=0.7	0.68	0.700	100.0
Commitment ϕ_1=1.0	0.16	1.000	100.0

Key finding: The Nash Trap pattern holds: selfish RL converges to $\bar{\lambda} \approx 0.49$ (Nash Trap) with 0% survival under 30% Byzantine agents. The commitment floor $\phi_1 \geq 0.7$ is required for survival in this environment—consistent with the Moral Commitment Spectrum.

U Constrained Policy Optimization: Learned Commitment Floor

To demonstrate that the commitment floor $\phi_1^* = 1.0$ emerges from principled optimization (not heuristic selection), we formulate and solve a constrained optimization problem via gradient-based Lagrangian duality:

$$\max_{\phi_1} \mathbb{E}\left[\sum_t r_t\right] \quad \text{s.t.} \quad \Pr(R_T > 0) \geq 1 - \delta \tag{19}$$

The Lagrangian relaxation $L(\phi_1, \mu) = \mathbb{E}[\sum r_t] - \mu \cdot (\delta - \hat{P}_{\text{surv}})$ is optimized via primal-dual gradient ascent: $\phi_1 \leftarrow \text{clip}(\phi_1 + \alpha_\phi \nabla_{\phi_1} L,\ 0,\ 1)$; $\mu \leftarrow \max(0,\ \mu + \alpha_\mu(\delta - \hat{P}_{\text{surv}}))$.

Table 27: Lagrangian dual CPO results (20 seeds, 200 outer iterations). All risk tolerances converge to $\phi_1^* = 1.0$.

Risk tolerance δ	**Learned ϕ_1^***	**Survival at convergence**
$\delta = 0.01$	1.000 ± 0.000	100%
$\delta = 0.05$	1.000 ± 0.000	100%
$\delta = 0.10$	1.000 ± 0.000	100%
$\delta = 0.20$	1.000 ± 0.000	100%
$\delta = 0.50$	1.000 ± 0.000	100%

Key finding: Across all five risk tolerance levels ($\delta \in [0.01, 0.50]$) and 20 random seeds, the Lagrangian dual CPO converges to $\phi_1^* = 1.000$ with zero variance. This confirms that unconditional commitment is the *unique safety-optimal solution*: any $\phi_1 < 1.0$ violates the survival constraint. The commitment floor is not a heuristic—it is the solution to a well-defined constrained optimization problem.

V Canonical QMIX Cross-Validation

To address the concern that baseline results may depend on implementation details, we re-implement QMIX following the original architecture [Rashid et al., 2018] and EPyMARL evaluation conventions [Papoudakis et al., 2021]: (i) hypernetwork-based mixing with absolute-value monotonicity constraints, (ii) per-agent Q-networks with shared parameters, (iii) ε-greedy exploration (not softmax), and (iv) target network updates every 200 steps.

Table 28: Canonical QMIX cross-validation (20 seeds). The EPyMARL-faithful implementation produces qualitatively identical Nash Trap convergence to our custom QMIX.

Implementation	$\bar{\lambda}$	**Survival %**	**Nash Trap?**
Custom QMIX (Table 18)	0.524 ± 0.022	66.5 ± 3.1	YES
Canonical QMIX (EPyMARL-style)	0.504 ± 0.144	68.6 ± 27.4	YES

Key finding: The canonical QMIX implementation converges to $\bar{\lambda} = 0.504$ (vs. 0.524 for our custom version) with comparable survival rates. Both implementations exhibit the Nash Trap: agents converge to individually rational cooperation levels that are insufficient for reliable collective survival ($\lambda_{\text{stab}} \geq 1.0$). The higher variance in the canonical implementation (± 0.144 vs. ± 0.022) reflects the sensitivity of ε-greedy exploration to seed initialization, but the *qualitative conclusion is invariant*: QMIX falls into the Nash Trap regardless of implementation details.