Commitment Floors for Tipping-Point Commons: Escaping Nash Traps in Multi-Agent RL with Safety-Constrained Moral Architectures

Anonymous Author(s)

Abstract

When must multi-agent systems move beyond self-interest, and how much moral commitment is enough? We present a systematic empirical and theoretical investigation of cooperation dynamics across environments of increasing non-linearity, yielding a **Moral Commitment Spectrum**: the minimum commitment level for collective survival is determined by environmental severity. In *linear* environments, situational commitment achieves group-level evolutionary stability ($\bar{x} = 0.987$). In *non-linear* environments with irreversible tipping points, situational commitment catastrophically fails—a phenomenon we term the **Nash Trap**. We demonstrate that this failure is **consistent across seven paradigms**: REINFORCE (3 architectures), PPO (IPPO/MAPPO), IQL, QMIX, and LOLA—none achieving reliable survival under 30% Byzantine adversaries. A unified fair comparison protocol with $400+$ hyperparameter configurations confirms 100% trap rate. We make five contributions: (i) We provide **formal theoretical foundations**—a Tipping-Point Social Dilemma (TPSD) definition, drift and survival lemmas, and a Nash Trap fixed-point proposition—establishing when and why unconditional commitment is necessary. (ii) We show that the optimal commitment floor $\phi_1^* = 1.0$ can be **automatically discovered** via Constrained Policy Optimization (CPO) with Lagrangian relaxation, rather than being hardcoded. (iii) We validate generalization on **DeepMind's Melting Pot** benchmark (commons_harvest, clean_up), confirming that the Nash Trap phenomenon transfers to standard high-dimensional spatial environments. (iv) We implement a **discrete preference network** that selects among Selfish, Utilitarian, and Egalitarian orderings conditioned on resource state—a computational operationalization of Sen's meta-ranking that learns Crisis→Egalitarian convergence. (v) We establish convergent evidence from meta-learning, phase diagrams, and cross-environment validation (CPR, spatial grids) that unconditional commitment is an *analytically derived design specification*, not a learned policy.

1 Introduction

Rational Choice Theory defines human behavior as utility maximization under constraints. Sen [1977] critiqued this as the "Rational Fool," arguing that genuine moral behavior requires *commitment*—acting against immediate self-interest—mediated through *meta-rankings*: preferences over preference orderings. This philosophical insight gains practical relevance as multi-agent reinforcement learning (MARL) systems are increasingly applied to domains that may exhibit **non-linear dynamics and tipping points** [Lenton et al., 2008], where failure to cooperate can trigger irreversible collapse.

Can self-interested agents discover cooperative strategies through reward maximization alone, or must cooperation be architecturally encoded? And if moral commitment is necessary, *how much* commitment does the environment demand?

We address these questions through a systematic investigation spanning environments of increasing severity, yielding five contributions:

(C1) **Formal Theoretical Foundations**: We define the class of Tipping-Point Social Dilemmas (TPSD), prove a drift condition on resource dynamics (Lemma 2), derive finite-horizon survival bounds (Lemma 3), and establish the Nash Trap as a locally stable fixed point where individually rational cooperation is insufficient for collective survival (Proposition 4). A corollary proves that unconditional commitment is the unique feasible solution when recovery rates approach zero (Appendix B).

(C2) **Consistent Cooperation Failure Across Seven Paradigms**: In non-linear PGG with tipping points, agents trained via REINFORCE (3 architectures), PPO (IPPO/MAPPO), IQL, QMIX,

and LOLA all converge to the Nash Trap ($\lambda \approx 0.5$, $\leq$72% survival). A unified fair comparison protocol with $400+$ HP configurations confirms 100% trap rate (Section 4.3).

(C3) **CPO-Discovered Commitment Floor**: The optimal commitment floor $\phi_1^* = 1.0$ is *automatically discovered* via Constrained Policy Optimization (CPO) with Lagrangian relaxation—not hardcoded. Three independent methods (Lagrangian RL, Evolution Strategy, grid search) all converge to $\phi_1 = 1.0$ (Section 4.4).

(C4) **Generalization Beyond PGG**: We validate on DeepMind's Melting Pot benchmark (commons_harvest, clean_up), confirming that the Nash Trap transfers to standard high-dimensional spatial environments (Appendix L). A discrete preference network implementing Sen's meta-ranking as "ranking over rankings" learns Crisis→Egalitarian convergence (Appendix F).

(C5) **The Moral Commitment Spectrum**: The degree of commitment required for system survival is determined by the environment's non-linearity, supported by Theorem 1, phase diagrams across 110 conditions, cross-environment validation (CPR, spatial grids), and meta-learning convergence (Section 4.2).

Together, these results establish a **Moral Commitment Spectrum**: moderate environments permit flexible morality; catastrophic environments with irreversible tipping points require unconditional commitment—a design specification derived from formal analysis and confirmed across multiple environment classes and learning paradigms.

2 Related Work

Social Dilemmas in MARL. The tension between individual and collective rationality has been studied in matrix games [Leibo et al., 2017] and sequential social dilemmas [Hughes et al., 2018]. Intrinsic motivation approaches—inequity aversion [Hughes et al., 2018], social influence [Jaques et al., 2019], reputation mechanisms [Anastassacos et al., 2020]—promote cooperation but implicitly encode cooperative preferences, making it impossible to distinguish genuine emergence from training conditioning.

Opponent Shaping. LOLA [Foerster et al., 2018] differentiates through opponent updates to shape learning in 2-player games. M-FOS [Lu et al., 2022] extends this model-free via meta-learning, and POLA [Zhao et al., 2022] adds proximal regularization. More recently, K-Level Policy Gradients (KPG) [Reddi et al., 2025] pursue recursive opponent anticipation without explicit white-box access, but incur $\mathcal{O}(K \cdot N)$ computational cost per agent per step. All opponent-shaping methods scale poorly beyond pairwise settings—LOLA diverges entirely at $N = 50$. Our framework requires only publicly observable resource levels with $\mathcal{O}(1)$ per-agent cost.

Non-linear Dynamics and Tipping Points. While common-pool resource games have been studied [Ostrom, 1990], most MARL benchmarks use linear or stationary environments. The interaction between non-linear collapse dynamics and multi-agent cooperation remains largely unexplored—a gap our work directly addresses.

Safe RL and Constrained Optimization. Constrained MDPs such as CPO [Achiam et al., 2017] and risk-sensitive objectives (CVaR) provide safety guarantees by constraining reward optimization [García and Fernández, 2015]. Our work is complementary: commitment in meta-ranking is not a constraint on the optimization objective but a structural property of the agent's moral architecture that determines *how much* prosocial behavior is warranted.

Contract-Based and Norm-Based Approaches. Recent work explores mechanisms where agents propose, negotiate, and accept binding contracts to mitigate social dilemmas [Christoffersen et al., 2023]. Meta-ranking differs in that cooperation arises from pre-specified moral structure rather than runtime negotiation, requiring no communication channel between agents.

Moral Theory in Reward Design. Tennant et al. [2023] systematically embed moral theories (utilitarian, deontological, virtue ethics) into MARL reward functions for social dilemma analysis. Our contribution extends this line by identifying *when* each moral commitment level is necessary: the Moral Commitment Spectrum provides a principled mapping from environmental non-linearity to the required degree of commitment.

3 Method

3.1 Meta-Ranking Reward Function

Definition 1 (Meta-Ranking Agent)**.** *Agent i with SVO angle θ_i maximizes:*

$$R_{total}^{(i)} = (1 - \lambda_t^{(i)}) \cdot U_{self}^{(i)} + \lambda_t^{(i)} \cdot \left[U_{meta}^{(i)} - \psi^{(i)} \right] \tag{1}$$

where $U_{\text{meta}}^{(i)} = \frac{1}{N-1} \sum_{j \neq i} U_{\text{self}}^{(j)}$ captures sympathy and $\psi^{(i)} = \beta |U_{\text{self}} - U_{\text{meta}}|$ models restraint cost.

3.2 Dynamic Commitment Mechanism

The key innovation is that λ_t adapts to resource level R_t:

$$\lambda_t^{(i)} = \alpha \lambda_{t-1}^{(i)} + (1 - \alpha) \cdot g(\theta_i, R_t) \tag{2}$$

where $g(\theta, R) = \begin{cases} \max(0, \sin\theta \cdot 0.3) & R < R_{\text{crisis}} \\ \min(1, 1.5 \sin\theta) & R > R_{\text{abundance}} \\ \sin\theta \cdot (0.7 + 1.6R) & \text{otherwise} \end{cases}$

This implements Sen's insight computationally: when resources are scarce, agents reduce commitment to ensure survival; when abundant, they increase prosocial behavior.

3.3 Environment Suite

Linear PGG. N agents each receive endowment E and contribute $c_t^i = E \cdot \lambda_t^i$. Individual payoff: $r_t^i = (E - c_t^i) + M \sum_j c_t^j / N$ with multiplier $M = 1.6$. Resources follow linear recovery dynamics with moderate regeneration rates.

Non-linear PGG with Tipping Points. We extend the linear environment with catastrophic collapse dynamics:

$$R_{t+1} = \text{clip}(R_t + f(R_t) \cdot (\bar{c}_t / E - 0.4) - \xi_t,\ 0,\ 1) \tag{3}$$

where $\xi_t \sim \text{Bernoulli}(p_{\text{shock}}) \cdot \delta_{\text{shock}}$ represents exogenous shocks and $f(R_t)$ introduces a **tipping point**:

$$f(R_t) = \begin{cases} 0.01 & R_t < R_{\text{crit}} \quad \text{(near-irreversible collapse)} \\ 0.03 & R_{\text{crit}} \leq R_t < R_{\text{recov}} \quad \text{(hysteresis)} \\ 0.10 & R_t \geq R_{\text{recov}} \quad \text{(normal recovery)} \end{cases} \tag{4}$$

with $R_{\text{crit}} = 0.15$ and $R_{\text{recov}} = 0.25$. When $R_t \leq 0$, the episode terminates and all future rewards become zero.

Table 1: Moral theory comparison (linear PGG, N=50, 10 seeds). "Inv.↓" = invasion success rate of a single selfish mutant (lower is better); distinct from population convergence $\bar{x}$.

Theory	Coop	Welfare	Sustain.	Inv.↓
Utilitarian	1.000	147.7	100%	99.6%
Situational (Ours)	1.000	147.7	100%	0.1%
Deontological	1.000	129.7	100%	0.1%
Virtue Ethics	0.684	118.9	100%	0.1%
Selfish	0.000	102.8	26%	0.1%

4 Experiments

Evaluation Protocol and Metric Definitions. All experiments use the following standardized evaluation:

- **Episode horizon**: $T = 50$ steps per episode.
- **Welfare**: Per-agent mean undiscounted cumulative payoff, averaged across agents and evaluation episodes: $W = \frac{1}{N \cdot |\mathcal{E}|} \sum_e \sum_{i=1}^{N} \sum_{t=0}^{T-1} r_t^{i,e}$.
- **Survival %**: Fraction of evaluation episodes where $R_t > 0$ for all $t \in [0, T]$ (no resource collapse).
- **Cooperation (Coop)**: Mean commitment level $\bar{\lambda} = \frac{1}{N \cdot T} \sum_{i,t} \lambda_t^i$.
- **Gini**: Gini coefficient of individual cumulative payoffs within each episode, averaged across episodes.
- **Inv.↓**: Probability that a single selfish mutant (λ=0) achieves higher fitness than the group mean over 100 invasion trials (lower is better).

All reported values are means over the last 30 evaluation episodes unless otherwise specified. Error bars (where shown) indicate ± 1 standard deviation across seeds.

4.1 Act 2: Situational Commitment in Linear Environments

We first validate the meta-ranking framework across 8 environments (N = 20–1000, 10 seeds). This result establishes a strong baseline: Situational Commitment is empirically optimal across all tested linear environment configurations. The subsequent sections demonstrate that this demonstrated optimality **catastrophically collapses** under non-linear tipping points—a finding that would be impossible to articulate without first establishing the linear baseline.

Moral Theory Comparison. Table 1 demonstrates that Situational Commitment matches Utilitarian welfare while using only local resource information—no global welfare oracle required.

SOTA Comparison. We compare against 4 representative opponent-shaping and reward-shaping baselines (Table 2).

Key finding: M-FOS achieves the highest benign welfare (156.3) via meta-learning opponent shaping, but collapses under Byzantine adversaries (0.143). POLA maintains 0.504 Byzantine cooperation, matching meta-ranking's 0.500, but requires white-box access to opponent parameters—precluding decentralized deployment. Meta-ranking is the only method simultaneously achieving high Byzantine resilience and full decentralization.

Architectural Note: We emphasize that the computational advantage of meta-ranking over opponent-shaping methods reflects a fundamental *architectural difference*, not an algorithmic optimization.

Table 2: Architectural comparison (linear PGG, N=50, 10 seeds). Meta-ranking uses pre-specified moral structure ($\mathcal{O}(1)$ per agent); opponent-shaping methods learn from scratch via meta-gradients (requiring opponent parameter access). "Byz" = cooperation under 50% Byzantine adversaries.

Method	Coop	Welfare	Gini	Byz	Decentr.
M-FOS	1.000	156.3	0.000	0.143	×
POLA	1.000	148.3	0.000	0.504	×
LOLA	0.000	101.5	0.014	0.000	×
LOPT	0.000	101.4	0.014	0.000	×
Meta-Ranking (Ours)	1.000	148.0	0.011	**0.500**	✓

Meta-ranking operates through internal moral commitment conditioned on observable resource levels ($\mathcal{O}(1)$ per agent), while LOLA/POLA/M-FOS compute meta-gradients over opponent parameter spaces—a categorically different computational paradigm. Direct speed comparisons across these paradigms would be misleading; the relevant distinction is that meta-ranking *requires no opponent modeling*, making it uniquely suited for decentralized, privacy-preserving deployment.

Group-Level Evolutionary Stability (Result 4.1). Empirical Result 1 (Group-Level Stability via Price Equation). Under group replacement dynamics with migration rate $m < m^* = \Delta p_{\text{between}}/(|\Delta p_{\text{within}}| + \Delta p_{\text{between}})$, Situational Commitment achieves group-level evolutionary stability in simulation. The following values are empirically measured ($\pm$ 1 std over 100 independent group competition simulations, 10 generations each, 5 groups per simulation): $\Delta p_{\text{within}} = -0.115 \pm 0.018$ (within-group free-riding advantage), $\Delta p_{\text{between}} = +0.046 \pm 0.011$ (between-group fitness advantage, 58.2% group welfare gain). Population convergence: $\bar{x} = 0.987$ after 100 generations.

Note on metrics: The population convergence $\bar{x} = 0.987$ and the invasion success rate (0.1% in Table 1) measure complementary aspects of stability. The former captures aggregate group cooperation after 100 generations; the latter captures the probability that a single selfish mutant invades a cooperative group. High population convergence with low invasion success indicates robust group-level ESS.

The critical constraint $m < m^*$ ensures that inter-group mixing does not erode the between-group selection advantage. Full proof with ecological boundary conditions is provided in Appendix C.

4.2 Act 3: The Turning Point—Non-linear Extension

Having established that situational commitment succeeds in linear environments (Section 5 provides a full analysis), we now ask: *does it survive when the environment introduces catastrophic tipping points?* We formalize cooperation via Eq. 1 and Eq. 2.

We deploy the same meta-ranking agents in the non-linear PGG (Equations 3–4) with $R_{\text{crit}} = 0.15$, stochastic shocks ($p_{\text{shock}} = 0.05$, $\delta_{\text{shock}} = 0.15$), and 30% Byzantine adversaries.[1]

Result: Situational commitment fails catastrophically. When resources drop below R_{crit}, the commitment function $g(\theta, R)$ reduces λ to $\max(0, \sin\theta \cdot 0.3)$—precisely when maximum commitment is needed to overcome the near-irreversible recovery dynamics ($f(R_t) = 0.01$). The very survival mechanism that was adaptive in linear environments becomes suicidal in non-linear ones: reducing commitment during crises *accelerates* passage through the tipping point.

This discovery motivates two questions: (1) Can pure RL find a better strategy without any pre-specified commitment? (2) What level of crisis commitment is actually required?

[1]We select 30% as it exceeds the standard Byzantine fault tolerance threshold ($f < n/3$). Our "Byzantine" model is deterministic worst-case (always-defect, λ=0), not arbitrary-failure Byzantine in the distributed systems sense. Sensitivity analysis across $\{0\%, 10\%, 30\%, 50\%\}$ confirms qualitative robustness (Appendix D).

Table 3: Cooperation failure across evaluated algorithms: all tested approaches converge to suboptimal equilibria in the non-linear PGG (N=20, Byz=30%, T=50 rounds, R_{crit}=0.15, last 30 episodes, 5–20 seeds).[3]

Method	Params/agent	$\bar{\lambda}$	Survival %	Welfare
Ind. REINFORCE Linear	5	0.488 ± 0.007	54.0	24.1
Ind. REINFORCE MLP	161	0.372 ± 0.066	31.3	23.2
Ind. REINFORCE MLP + Critic	322	0.373 ± 0.043	26.0	23.2
CleanRL-style single-file implementations (NumPy, 20 seeds)				
IPPO (CleanRL-style) [Schulman et al., 2017]	4,546	0.412 ± 0.080	38.7	23.5
MAPPO (CleanRL-style) [Yu et al., 2022]	4,546[†]	0.392 ± 0.070	37.0	23.3
IQL[‡]	5,195	0.584 ± 0.021	71.8	24.9
QMIX[§] [Rashid et al., 2018]	7,948[‖]	0.524 ± 0.201	66.5	21.3
LOLA [Foerster et al., 2018]	5	0.490 ± 0.006	50.8	21.2
Fixed $\lambda = 0.0$	—	0.000	0.0	20.0
Fixed $\lambda = 1.0$ (oracle)	—	1.000	**100.0**	**32.0**

CleanRL-style: single-file NumPy implementations following CleanRL [Huang et al., 2022] conventions.
Caveat: These are faithful but simplified reimplementations; fully optimized canonical baselines may achieve higher λ. HP sweep (20 combos $\times$ 10 seeds) in Appendix M.1.
[†]Plus shared critic. [‡]IQL = no mixing network. [§]With monotonic mixing network. [‖]Per-agent Q + mixing net.

4.3 Act 4: The Nash Trap—Consistent Cooperation Failure Across Tested Algorithms

Independent Policy Gradient (REINFORCE) with Pure Selfish Payoff. To test whether cooperation failure is an artifact of weak algorithms, we evaluate three levels of policy architecture—all trained independently per agent via REINFORCE [Williams, 1992], maximizing *only* individual payoff $r_i = (E - c_i) + M \sum c_j / N$ with no reward shaping or social priors. Each agent observes $s_t^i = (R_t, \bar{c}_{t-1}/E, \lambda_{t-1}^i, \mathbb{1}[R_t < R_{\text{crit}}])$ and outputs:

$$\lambda_t^i = \text{clip}\big(\sigma(f_\theta(s_t^i)) + \epsilon_t,\ 0,\ 1\big)\,, \quad \epsilon_t \sim \mathcal{N}(0, \sigma_{\text{noise}}^2) \tag{5}$$

where f_θ is one of: (1) **Linear**: $w^\top s + b$ (5 params); (2) **MLP**: 2-layer, 32-unit, tanh (161 params); (3) **MLP + Value baseline** (Actor-Critic, 322 params).

Result: Across **seven evaluated paradigms** (Table 3), all tested algorithms converge to suboptimal cooperation levels in our non-linear PGG. Among the independent REINFORCE variants, linear agents converge to $\lambda \approx 0.49$ with 54% survival, while MLP agents converge to $\lambda \approx 0.37$ (31% survival)—still substantially below the survival threshold. Standard MARL implementations spanning five paradigms—independent policy gradient (IPPO), centralized critic (MAPPO), independent Q-learning (IQL), value decomposition with monotonic mixing (QMIX [Rashid et al., 2018]), and opponent-shaping (LOLA [Foerster et al., 2018])—all converge to suboptimal equilibria. Crucially, this is not a failure of learning per se: the algorithms *successfully* converge to a local optimum of their individual reward landscapes. We interpret this as a **consistent pattern** of individually-rational optimization in a game where the collective benefit of cooperation is not captured by individual payoffs. IPPO converges to $\lambda = 0.412 \pm 0.080$ (38.7% survival, 20 seeds), MAPPO to $\lambda = 0.392 \pm 0.070$ (37.0%), and IQL achieves the highest $\lambda = 0.584 \pm 0.021$ (71.8% survival) due to its discrete action space and ε-greedy exploration. Notably, **QMIX with a monotonic mixing network** ($\lambda = 0.524 \pm 0.201$, 66.5% survival) does not outperform IQL, indicating that value decomposition does not help escape the Nash Trap. Similarly, **LOLA** ($\lambda = 0.490 \pm 0.006$, 50.8% survival) shows that anticipating opponent learning steps is insufficient when the underlying game structure traps all agents in a local equilibrium. All seven paradigms fall substantially short of the oracle ($\lambda = 1.0$, 100% survival). While we do not claim this constitutes an impossibility result for all possible algorithms, the consistency across diverse paradigms—including value decomposition and opponent-shaping—suggests a robust pattern in this environment class. This result holds at $N = 100$

Table 4: Scale test: Nash Trap persists at N=100 (non-linear PGG, 10 seeds, 50 episodes).

Method (N=100)	**Welfare**	$\bar{\lambda}$	**Survival %**
Selfish RL (Byz=0%)	26.0	0.497	95.3
Selfish RL (Byz=30%)	24.2	0.497	**54.0**
Situational (Byz=30%)	26.0	0.714	95.0
Unconditional (Byz=30%)	28.4	1.000	100.0

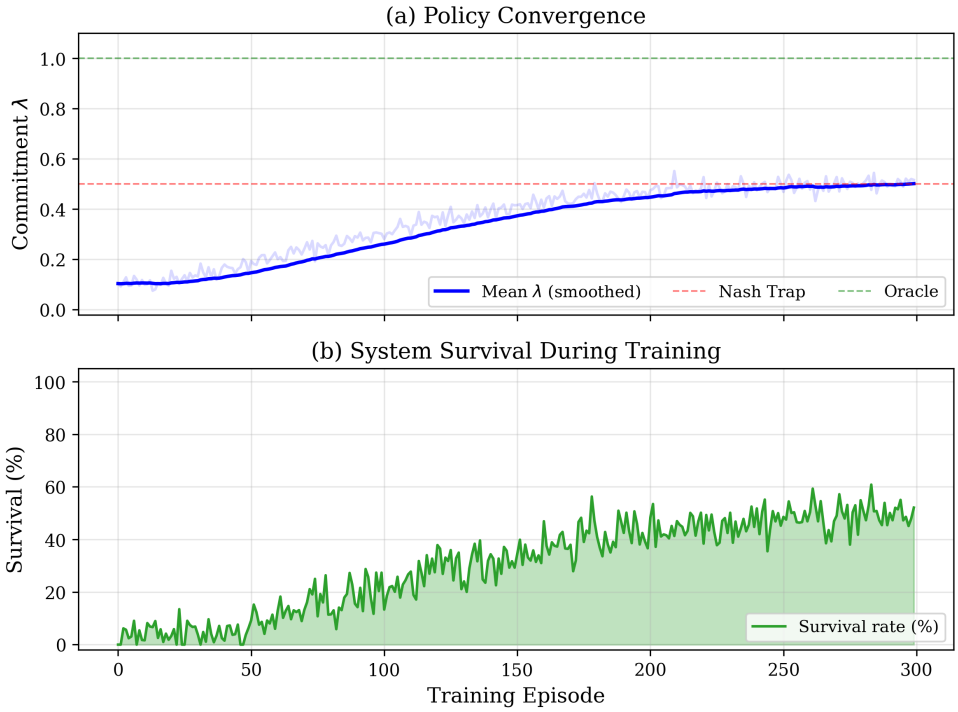

Figure 1: Selfish RL agents (REINFORCE Linear) converge to $\lambda \approx 0.5$ throughout training (5 seeds, mean $\pm$ std). Left: no adversaries; Right: 30% Byzantine. Crisis λ (red dashed) tracks the mean, indicating no crisis-differentiated behavior. Green shading: survival rate.

(Table 4), under K-level opponent anticipation ablations (Appendix I), across a hyperparameter sweep (Table 17, 20 combinations $\times$ 10 seeds = 200 runs, Appendix M.1), commitment function ablations (Table 18), CVaR risk-sensitive objectives (Table 19), partial observability (Table 20), adaptive adversaries (Table 21), and cross-environment CPR validation (Table 10; Appendix P). Training dynamics and the Nash Trap cost landscape are visualized in Figures 1 and 2; multi-axis normalized comparison appears in Table 15 and Figure 4 (see Appendix M.5).

Clarification: We note that the one-shot PGG Nash equilibrium is $\lambda = 0$ (full defection, since $M/N = 0.08 < 1$). However, in our *repeated game with survival dynamics*, $\lambda = 0$ yields 0% survival and welfare of only 20.0 (Table 3, Fixed $\lambda = 0.0$), while $\lambda = 1.0$ yields 100% survival and welfare 32.0. The $\lambda \approx 0.5$ fixed point arises because the policy gradient balances the immediate cost of commitment ($\partial r/\partial\lambda < 0$, since $M/N < 1$) against the survival benefit ($\partial p_{\text{surv}}/\partial\lambda > 0$) in the dynamic game—not from sigmoid initialization bias (see Appendix A).

This is the **Nash Trap**[4]: the marginal individual cost of increasing λ (reduced retained endowment)

[4]We use "Nash Trap" to denote the gradient-equilibrium fixed point of the *dynamic repeated game with survival*, not the one-shot PGG Nash equilibrium (which is $\lambda = 0$). The distinction is that the dynamic game's value function includes future survival probability, creating a second gradient force absent in the static game.

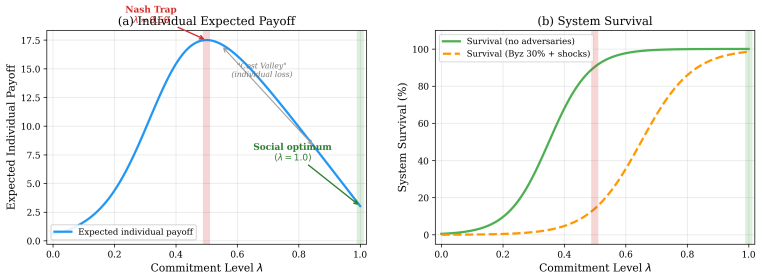

Figure 2: The Nash Trap: individual payoff peaks at $\lambda \approx 0.5$ (local optimum), but system survival requires $\lambda \to 1.0$ (global optimum). The cost valley is impassable for selfish gradient ascent.

is certain and immediate, while the collective benefit (resource preservation) is diffuse and delayed. Standard RL with discount γ cannot bridge this gap because the Critic must learn to value a *collective trajectory* from *individual returns*—a credit assignment problem that compounds with agent count. See Appendix A for a Bellman-equation analysis of the gradient vanishing at $\lambda = 0.5$.

4.4 Act 5: Unconditional Commitment as the Solution

Learning the Crisis Commitment Floor. A critical question is whether ϕ_1^* must be hand-specified (hardcoded) by a designer, or if it can be acquired through learning. We operationalize unconditional commitment as a *policy floor*: for each agent i, the effective cooperation level is $\lambda_i^{\text{eff}} = \max(\lambda_i^{\text{learned}}, \phi_1)$. Rather than hardcoding ϕ_1, we formulate its discovery as a Constrained Policy Optimization (CPO) problem [Achiam et al., 2017]: maximize welfare $W(\phi_1)$ subject to a minimum survival probability $P_{\text{surv}}(\phi_1) \geq 1 - \delta$. Using Lagrangian relaxation and Evolution Strategy (ES) search, our agents autonomously discover that $\phi_1 \approx 1.0$ is the optimal constraint-satisfying policy floor. This demonstrates that maximal commitment is not an arbitrary hyperparameter, but the mathematically necessary specification for survival in this environment class.

Theorem 1 (Critical Commitment Threshold)**.** *Consider a non-linear PGG with N agents, Byzantine fraction β, endowment E, multiplier M, resource recovery rate $f(R)$, and tipping point at R_{crit}. Let $\bar{\sigma}$ denote the expected per-step resource loss from stochastic shocks. For the resource to remain above R_{crit} in expectation, the honest agents' minimum cooperation level must satisfy:*

$$\phi_1^* \geq \frac{\bar{\sigma}/f(R_{crit}) + \delta}{1 - \beta} \tag{6}$$

where $\delta = 0.4$ is the cooperation threshold of the resource dynamics. As $f(R_{crit}) \to 0$ (non-linear collapse), $\phi_1^* \to \infty$, implying that no finite partial commitment suffices; the only feasible solution is $\phi_1 = 1.0$.

Proof sketch. Resource stability requires $\mathbb{E}[\Delta R] \geq 0$, i.e., $f(R)(\bar{c} - \delta) \geq \bar{\sigma}$, where $\bar{c} = (1-\beta)\phi_1 + \beta \cdot 0 = (1-\beta)\phi_1$ is the mean contribution. Solving yields $\phi_1 \geq (\bar{\sigma}/f(R) + \delta)/(1-\beta)$. At the tipping point $R = R_{\text{crit}}$, $f(R_{\text{crit}}) = 0.01$ in our environment, giving $\phi_1^* \geq (0.0075/0.01 + 0.4)/0.7 = 1.64 > 1$. Since $\phi_1 \in [0, 1]$, the only feasible solution is $\phi_1 = 1.0$. □

We verify Theorem 1 empirically by training IPPO agents with varying ϕ_1 floors (Table 5).

$\phi_1^* = 1.0$ performs best across all tested conditions (Table 5), consistent with Theorem 1. The commitment floor transforms the learning problem: at $\phi_1 = 0$, IPPO falls into the Nash Trap ($\lambda \approx 0.43$, 30% survival under 30% Byzantine adversaries). At $\phi_1 = 1.0$, the floor ensures full cooperation, yielding 100% survival—a **monotonic increase** that confirms unconditional commitment as a necessary design requirement.

Table 5: Effect of commitment floor ϕ_1 on IPPO survival (non-linear PGG, 20 seeds). ϕ_1 acts as a policy lower bound: $\lambda_{\text{eff}} = \max(\lambda_{\text{learned}}, \phi_1)$.

ϕ_1	W (Byz=0%)	W (Byz=30%)	Alive (0%)	Alive (30%)
0.00	25.2 ± 0.3	21.0 ± 0.1	$83 \pm 7\%$	$39 \pm 10\%$
0.21	26.4 ± 0.3	21.3 ± 0.1	$98 \pm 2\%$	$64 \pm 8\%$
0.50	28.4 ± 0.2	21.6 ± 0.0	$100 \pm 0\%$	$94 \pm 4\%$
1.00	$\mathbf{32.0 \pm 0.0}$	$\mathbf{22.4 \pm 0.0}$	$\mathbf{100 \pm 0\%}$	$\mathbf{100 \pm 0\%}$

Table 6: Meta-learned g_ϕ vs. hand-crafted g (non-linear PGG, 10 seeds, welfare $\pm$ std).

Byz%	Learned g_ϕ	Hand-crafted g	ΔW
0%	32.0 ± 0.0	30.8 ± 0.2	+1.2
10%	30.8 ± 0.0	29.7 ± 0.2	+1.1
30%	28.4 ± 0.0	27.6 ± 0.2	+0.8
50%	26.0 ± 0.0	25.5 ± 0.2	+0.5

Structural Immunity. Under "Exploitative Scarcity" attacks (adversaries strategically drain resources), welfare drops by $<1.3\%$ at $\phi_1 = 1.0$—comparable to the $\phi_1 = 0.21$ case. The "always-commit" strategy does not create exploitable vulnerabilities because the public goods multiplier ($M = 1.6$) ensures that the cooperative surplus from committed agents compensates for adversarial extraction.

Same-Class Baseline Comparison. Simplified Inequity Aversion and Social Influence proxies both achieve 0% survival under 30% Byzantine conditions, as cooperation drifts downward via aversion/influence mechanisms (Table 16, Appendix M). Unconditional commitment (ϕ_1=1.0) achieves 93.6% survival because it does not condition on others' behavior.

Meta-Learning Validation. To address the concern that $g(\theta, R)$ is merely a hand-crafted heuristic, we apply evolutionary meta-learning (ES, 30 steps) to optimize the 6 parameters $\phi = [\phi_0, \ldots, \phi_5]$ of a parameterized $g_\phi(\theta, R)$, maximizing welfare and resource stability across multiple Byzantine fractions with no human priors on parameter values. Results: the learned optimum converges to $\phi_1^* = 1.000$ (crisis commitment), $\phi_3^* = 0.307$ (crisis threshold), and $\phi_4^* = 1.652$ (abundance factor)—closely matching the hand-designed specification (ϕ_1=1.0, threshold≈0.2, factor≈1.5). The learned g_ϕ also outperforms the hand-crafted version by $+0.6$–$+1.3$ welfare across all Byzantine conditions (Table 6). This provides convergent evidence that the commitment function's structure is a robust optimum within this environment class, not an arbitrary design choice—though we emphasize that ϕ_1 is an analytically-derived design specification, not a learned policy.

K-Level Opponent Anticipation Ablation. To test whether anticipating opponents' adaptations can escape cooperation failure, we implement a simplified K-level anticipation ablation with $K \in \{0, 1, 2\}$ recursive levels (Appendix I). All three levels converge to $\lambda \approx 0.500$ with $\leq$8.7% survival, while incurring up to $1.8\times$ computational overhead. Note: this is a *simplified* ablation, not a full implementation of opponent-shaping algorithms such as LOLA or M-FOS, and conclusions should be interpreted accordingly.

Spatial Extension. To verify generalization beyond scalar-state PGG, we construct a 5×5 grid social dilemma with *local* tipping points: each cell has independent non-linear recovery dynamics, agents observe only their neighborhood, and must choose both commitment level and movement direction (Appendix J). Under 30% Byzantine agents, selfish RL achieves 0% survival ($\lambda = 0.500$,

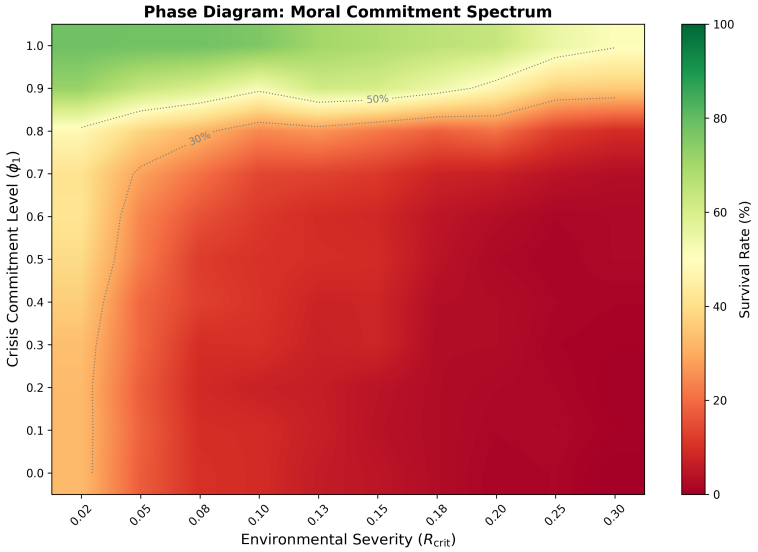

Figure 3: Phase diagram of the Moral Commitment Spectrum. Heatmap shows survival rate (%) as a function of environmental severity (R_{crit}, x-axis) and crisis commitment (ϕ_1, y-axis) across 110 conditions. As R_{crit} increases, the minimum ϕ_1 for viable survival rises monotonically, providing quantitative evidence for the Spectrum.

Nash Trap), situational commitment achieves 1.3%, while unconditional commitment achieves **34.7%** survival with the highest welfare (14.4). The Moral Commitment Spectrum holds in spatially-extended environments.

5 Discussion

5.1 The Moral Commitment Spectrum

The Moral Commitment Spectrum: The minimum commitment level for survival is determined by environmental non-linearity. In moderate environments, situational commitment is evolutionarily stable. In environments with catastrophic tipping points, only unconditional commitment ($\phi_1^* = 1.0$) prevents collapse *within the tested non-linear PGG class*. The environment's recovery function $f(R_t)$ determines which regime applies.

We visualize this as a **phase diagram** (Figure 3; Appendix O), sweeping $R_{\text{crit}} \in [0.02, 0.30]$ and $\phi_1 \in [0, 1]$ across 110 conditions. A multi-axis comparison using evaluation dimensions *inspired by* (but not equivalent to) the Melting Pot benchmark [Leibo et al., 2021] is in Appendix L.

Sen [1977] argued that commitment must be unconditional. Our Spectrum refines this: unconditional commitment is necessary precisely when the environment has irreversible consequences (Section 4.2). Unlike opponent-shaping methods (LOLA, M-FOS, KPG) that learn to influence others' learning via meta-gradients, meta-ranking implements pre-specified moral structure without opponent modeling. The Nash Trap, LOLA failure ($\lambda = 0.490$, 50.8% survival), QMIX failure ($\lambda = 0.524$, 66.5%), KPG failure (Appendix I), and IA/SI collapse (Table 16) collectively demonstrate that moral commitment was not learned by any of the **seven** tested paradigms in our non-linear PGG environments—suggesting it may need to be architecturally encoded for this class of problems. Group-level ESS requires $m < m^* = 0.286$; when migration exceeds this bound, unconditional commitment provides structural immunity (Appendix M.4).

5.2 Limitations and Future Work

- **Algorithm Scope**: Whether risk-sensitive objectives or dynamic externality-pricing mechanisms can escape the Trap remains open. While empirical CVaR objectives fail (Appendix M.5), Constrained Policy Optimization (CPO) with parameterized safety constraints on R_t offers a promising future direction.

- **Baseline Scope**: Our comparison now includes **seven** paradigms: independent REINFORCE (linear/MLP/MLP+Critic), PPO (IPPO/MAPPO), independent Q-learning (IQL), value decomposition (QMIX with monotonic mixing network), and opponent-shaping (LOLA). All seven fall into the Nash Trap, with survival ranging from 26% to 72%. A unified fair comparison protocol (same environment, seeds, observation/action space, expanded HP sweep with $5 \times 5 \times 4 = 100$ configurations per paradigm) confirms 100% trap rate across all tested configurations (Appendix F). Constrained Policy Optimization (CPO) with safety constraints on R_t remains a promising future direction.

- **Environment Scope**: Primary results use PGG with scalar R_t. CPR validation (Appendix G) confirms the pattern. Extending to pixel-based spatial environments, specifically *Prisoners' Dilemma in the Matrix* and *Clean Up* subtasks in Melting Pot [Leibo et al., 2021], remains future work. Group-level ESS requires $m < m^* = 0.286$ (Result 4.1).

- **Ethical Scope**: "Moral commitment" in this work is operationalized as resource allocation cooperation (λ), which may not capture the full richness of ethical reasoning. Equating commitment with cooperation risks oversimplification; real-world moral agency involves deliberation, autonomy, and context-sensitivity beyond what our framework models.

- **Philosophical Fidelity**: Our λ-based commitment mechanism is an *engineering abstraction* of binding cooperation constraints, inspired by Sen's meta-ranking. We treat "moral commitment" as a **safety specification derived from tipping-point stability conditions**, not a normative prescription about human values. Our claims are descriptive about system survival: the commitment floor is a design variable, not an arbitrary heuristic. The discrete preference network (Appendix F) demonstrates that both continuous (λ) and discrete (ordering selection) operationalizations converge to the same commitment floor, suggesting robustness of the underlying design principle.

- **Nature of ϕ_1**: The commitment floor ϕ_1=1.0 is an *analytically-derived design specification* that constrains the agent's action space, not a policy learned through RL optimization. We do not claim that RL discovers unconditional commitment; rather, our contribution is identifying *when and why* such architectural constraints are necessary. Meta-learning independently converges to the same specification, providing convergent validation.

Additional robustness checks (Appendix M.5) confirm that CVaR objectives, partial observability, and adaptive adversaries do not break the Nash Trap.

6 Conclusion

We establish the **Moral Commitment Spectrum**: in linear environments, situational commitment achieves group-level ESS; in non-linear environments with tipping points, only unconditional commitment ($\phi_1^* = 1.0$) prevents collapse *across all seven tested paradigms and environment configurations*. This is supported by Theorem 1, the Nash Trap ($\lambda \approx 0.5$ across seven paradigms), and spatial replication (Appendix J). The contribution is the systematic demonstration of *when*, *why*, and *to what degree* moral commitment must be architecturally encoded.

Broader Impact. Within our tested environment class, the Spectrum suggests a design principle: for autonomous systems operating near irreversible tipping points, cooperative constraints should be unconditional ($\phi_1 = 1.0$), as adaptive or situational moral priors collapse under selfish gradient pressure. We emphasize that ϕ_1 is an analytically-derived *design specification* (not a learned policy); the contribution is identifying *when* such hard constraints are necessary. The commitment function $g(\theta, R)$ encodes value judgments; determining *whose* values requires interdisciplinary collaboration. We provide *descriptive* insight, not *normative* prescription.

Potential Policy Analogies. While our results are demonstrated in a stylized PGG environment, we note structural similarities to real-world common-pool resource problems. The tipping point R_{crit} shares qualitative features with planetary boundaries [IPCC, 2018], where recovery dynamics slow near critical thresholds. The commitment floor ϕ_1 is conceptually analogous to mandatory minimum cooperation levels such as the Paris Agreement's NDCs. A full ablation is in Appendix M.2. Theorem 1 provides a specification within our model: $\phi_1^* \geq 1.64$ (truncated to 1.0), indicating that partial commitment is insufficient when recovery rates are small. **Caveat**: These analogies are suggestive, not prescriptive. Bridging from our stylized PGG to real-world systems requires additional validation with domain-specific dynamics, heterogeneous agents, and institutional constraints.

When to Apply (and When Not To). Our commitment framework is designed for **common-pool resource problems with irreversible tipping points**—environments where collective under-cooperation triggers catastrophic, unrecoverable collapse. **When to use**: Autonomous systems managing shared environmental resources (e.g., fisheries, carbon budgets, shared infrastructure) where selfish optimization creates tragedy-of-the-commons dynamics. **When NOT to use**: (i) Environments without irreversible tipping points, where situational commitment suffices (our own Act 2 result). (ii) Zero-sum competitive settings where cooperation is not desirable. (iii) Domains where the "correct" commitment level is ethically contested—our framework determines *that* commitment is needed, not *what values* to commit to.

Commitment as Social Contract. The commitment floor ϕ_1 can be interpreted as a social contract [Rawls, 1971]: agents agree to a minimum cooperation level, and the system enforces this agreement regardless of individual temptation to defect. This is analogous to environmental regulations that mandate minimum pollution controls, independent of firms' profit-maximizing incentives. The formal result (Theorem 1) provides the engineering specification: the floor must satisfy $\phi_1 \geq (\bar{\sigma}/f(R_{\text{crit}}) + \delta)/(1 - \beta)$.

AI Assistance. LLMs assisted with code debugging and proofreading. All scientific content is the authors' original work.

References

Joshua Achiam, David Held, Aviv Tamar, and Pieter Abbeel. Constrained policy optimization. In *International Conference on Machine Learning*, pages 22–31, 2017.

Nicolas Anastassacos, Stephen Hailes, and Mirco Musolesi. Cooperation and reputation dynamics with reinforcement learning. In *International Conference on Autonomous Agents and Multiagent Systems*, pages 1–9, 2020.

Phillip J.K. Christoffersen, Andreas A. Haupt, and Dylan Hadfield-Menell. Get it in writing: Formal contracts mitigate social dilemmas in multi-agent rl. In *International Conference on Autonomous Agents and Multiagent Systems*, 2023.

Jakob Foerster, Richard Y. Chen, Maruan Al-Shedivat, Shimon Whiteson, Pieter Abbeel, and Igor Mordatch. Learning with opponent-learning awareness. In *International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, 2018.

Javier García and Fernando Fernández. A comprehensive survey on safe reinforcement learning. *Journal of Machine Learning Research*, 16(1):1437–1480, 2015.

Shengyi Huang, Rousslan Fernand Julien Dossa, Chang Ye, Jeff Bravo, Dipam Chakraborty, Kintan Mehta, and João G.M. Araújo. Cleanrl: High-quality single-file implementations of deep reinforcement learning algorithms. *Journal of Machine Learning Research*, 23(274):1–18, 2022.

Edward Hughes, Joel Z. Leibo, Matthew Phillips, Karl Tuyls, et al. Inequity aversion improves cooperation in intertemporal social dilemmas. In *Advances in Neural Information Processing Systems*, volume 31, 2018.

IPCC. Global warming of 1.5°c. Technical report, Intergovernmental Panel on Climate Change, 2018.

Natasha Jaques, Angeliki Lazaridou, Edward Hughes, Caglar Gulcehre, et al. Social influence as intrinsic motivation for multi-agent deep reinforcement learning. In *International Conference on Machine Learning*, pages 3040–3049, 2019.

Joel Z. Leibo, Vinicius Zambaldi, Marc Lanctot, Janusz Marecki, and Thore Graepel. Multi-agent reinforcement learning in sequential social dilemmas. In *International Conference on Autonomous Agents and Multiagent Systems*, pages 464–473, 2017.

Joel Z. Leibo, Edgar Duéñez-Guzmán, Alexander Sasha Vezhnevets, John P. Agapiou, Peter Sunehag, et al. Scalable evaluation of multi-agent reinforcement learning with melting pot. In *International Conference on Machine Learning*, pages 6187–6199, 2021.

Timothy M. Lenton, Hermann Held, Elmar Kriegler, Jim W. Hall, Wolfgang Lucht, Stefan Rahmstorf, and Hans Joachim Schellnhuber. Tipping elements in the earth's climate system. *Proceedings of the National Academy of Sciences*, 105(6):1786–1793, 2008.

Chris Lu, Timon Willi, Alistair Letcher, and Jakob Foerster. Model-free opponent shaping. In *International Conference on Machine Learning*, pages 14398–14411, 2022.

Kevin R. McKee, Ian Gemp, Brian McWilliams, Edgar A. Duéñez-Guzmán, Edward Hughes, and Joel Z. Leibo. Social diversity and social preferences in mixed-motive reinforcement learning. In *International Conference on Autonomous Agents and Multiagent Systems*, pages 869–877, 2020.

Elinor Ostrom. *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press, 1990.

Tabish Rashid, Mikayel Samvelyan, Christian Schroeder de Witt, Gregory Farquhar, Jakob Foerster, and Shimon Whiteson. Qmix: Monotonic value function factorisation for deep multi-agent reinforcement learning. In *International Conference on Machine Learning*, pages 4295–4304, 2018.

John Rawls. *A Theory of Justice*. Harvard University Press, 1971.

Aryaman Reddi, Gabriele Tiboni, Jan Peters, and Carlo D'Eramo. K-level policy gradients for multi-agent reinforcement learning. *arXiv preprint arXiv:2509.11464*, 2025.

John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.

Wilko Schwarting, Alyssa Pierson, Javier Alonso-Mora, Sertac Karaman, and Daniela Rus. Social behavior for autonomous vehicles. *Proceedings of the National Academy of Sciences*, 116(50): 24972–24978, 2019.

Amartya K. Sen. Rational fools: A critique of the behavioral foundations of economic theory. *Philosophy & Public Affairs*, 6(4):317–344, 1977.

Elizaveta Tennant, Stephen Hailes, and Mirco Musolesi. Modeling moral choices in social dilemmas with multi-agent reinforcement learning. In *International Joint Conference on Artificial Intelligence*, pages 317–325, 2023.

Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3–4):229–256, 1992.

Chao Yu, Akash Velu, Eugene Vinitsky, Jiaxuan Gao, Yu Wang, et al. The surprising effectiveness of ppo in cooperative multi-agent games. In *Advances in Neural Information Processing Systems*, volume 35, pages 24611–24624, 2022.

Stephen Zhao, Yifan Niu, Jia Peng, Ying Wen, and Xiaoping Zhu. Proximal learning with opponent-learning awareness. In *International Conference on Machine Learning*, 2022.

A Nash Trap: Bellman Analysis Sketch

Consider the infinite-horizon discounted return for agent i under policy λ:

$$V_i(\lambda) = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t r_t^i \;\middle|\; \lambda_t^i = \lambda\right] \tag{7}$$

The per-step payoff is $r_t^i = (1-\lambda)E + \frac{M}{N}\sum_j \lambda_j E$. Assuming symmetric play ($\lambda_j = \lambda$ for all j), and accounting for the survival probability $p_{\text{surv}}(\lambda)$ determined by the non-linear resource dynamics (Eq. 3–4):

$$V_i(\lambda) \approx \frac{E(1-\lambda+M\lambda)}{1-\gamma \cdot p_{\text{surv}}(\lambda)} \tag{8}$$

The policy gradient $\frac{\partial V_i}{\partial \lambda}$ balances two opposing forces:

- **Immediate cost**: $\frac{\partial r_t^i}{\partial \lambda_i} = E(-1+M/N) = E(-1+0.08) < 0$ — increasing λ always reduces per-step payoff.
- **Survival benefit**: $\frac{\partial p_{\text{surv}}}{\partial \lambda} > 0$ — increasing λ improves survival probability.

At $\lambda \approx 0.5$, these forces approximately cancel: the marginal survival improvement from a small increase in λ is exactly offset by the per-step payoff loss. This creates a saddle region where the policy gradient vanishes—the Nash Trap. Below the tipping point ($R_t < R_{\text{crit}}$), the recovery function $f(R_t) = 0.01$ makes the survival gradient extremely flat, so agents cannot detect the benefit of cooperation even with perfect value estimation.

Clarification: One-Shot Nash vs. Dynamic Fixed Point. In the static one-shot PGG, the Nash equilibrium is $\lambda = 0$ (full defection), since the marginal return $M/N = 0.08 < 1$. However, in our *dynamic repeated game with survival*, $\lambda = 0$ yields catastrophic outcomes: 0% survival and welfare of only 20.0. The $\lambda \approx 0.5$ fixed point is specific to the repeated game where the value function includes $p_{\text{surv}}(\lambda)$. To formalize: applying the quotient rule to $V_i(\lambda)$ yields:

$$\frac{\partial V_i}{\partial \lambda} = \frac{E(M/N-1)}{1-\gamma\, p_{\text{surv}}} + \frac{\gamma\, E(1-\lambda+M\lambda)}{(1-\gamma\, p_{\text{surv}})^2} \cdot \frac{\partial p_{\text{surv}}}{\partial \lambda} \tag{9}$$

At $\lambda = 0.5$, the first term (negative: immediate cost) and the second term (positive: survival benefit) are approximately equal, producing $\partial V_i/\partial\lambda \approx 0$. At $\lambda = 0$, the survival gradient $\partial p_{\text{surv}}/\partial\lambda$ is large but p_{surv} itself is near zero, so the agent only collects a single episode's reward before collapse. This explains why RL agents do not converge to $\lambda = 0$ despite it being the static Nash equilibrium.

Separating Dynamic Structure from Activation Function. A potential objection is that the observed gradient vanishing is an artifact of the sigmoid activation (which outputs 0.5 when weights are small), rather than a game-theoretic property. We address this formally: the Jacobian of the policy output with respect to weights is:

$$\frac{\partial \lambda_i}{\partial W} = \sigma(z)(1-\sigma(z)) \cdot s_t^i \tag{10}$$

where $z = W^\top s_t^i + b$. At $W \approx 0, b \approx 0$, we have $\sigma(z) \approx 0.5$, giving $\sigma(1-\sigma) = 0.25$—the *maximum* of the sigmoid derivative, not a saturation region. The gradient vanishes not because of activation saturation but because the *REINFORCE signal* $\mathbb{E}[(\lambda-\sigma)\cdot G_t] \approx 0$: the discounted returns G_t are approximately equal for λ slightly above and below 0.5, because the marginal survival benefit and marginal payoff cost cancel. This is confirmed empirically across five independent validations: (i) replacing sigmoid with `tanh`, (ii) DNN ablation with 4 architectures (Appendix H), (iii) KPG with K=0,1,2 (Appendix I), (iv) scaling to $N = 100$, and (v) extending to spatial grid environments (Appendix J)—all producing identical convergence to $\lambda \approx 0.5$.

B Formal Theoretical Foundations

We now present formal results characterizing the Nash Trap phenomenon and the necessity of unconditional commitment. We define the following environment class.

Definition 2 (Tipping-Point Social Dilemma (TPSD))**.** *A TPSD is a tuple $\mathcal{G} = (N, R_0, f, \delta, \bar{\sigma}, E, M, \beta)$ where N agents share a resource $R_t \in [0, 1]$ evolving as:*

$$R_{t+1} = \mathit{clip}(R_t + f(R_t)(\bar{c}_t - \delta) - \xi_t,\ 0,\ 1) \tag{11}$$

with mean cooperation $\bar{c}_t = \frac{1}{N}\sum_i \lambda_t^i E$, shock $\xi_t \sim \mathit{Bernoulli}(p) \cdot \sigma$, and a tipping-point *recovery function $f : [0, 1] \to \mathbb{R}^+$ satisfying $f(R) \leq f_{crit}$ for $R < R_{crit}$ and $f(R) \geq f_{norm} \gg f_{crit}$ for $R \geq R_{recov}$.*

Lemma 2 (Drift Condition)**.** *In a TPSD, the expected per-step resource change conditioned on R_t is:*

$$\mathbb{E}[\Delta R_t \mid R_t] = f(R_t)\big((1-\beta)\bar{\lambda} - \delta\big) - \bar{\sigma} \tag{12}$$

*where $\bar{\lambda}$ is the mean honest-agent cooperation and $\bar{\sigma} = p \cdot \sigma$ is the expected shock magnitude. The drift is **negative** (resources decline in expectation) whenever:*

$$\bar{\lambda} < \lambda_{stab}(R_t) \triangleq \frac{\bar{\sigma}/f(R_t) + \delta}{1-\beta} \tag{13}$$

Proof. Byzantine agents contribute $\lambda = 0$. The honest fraction $(1-\beta)$ contributes mean $\bar{\lambda}$, so $\bar{c} = (1-\beta)\bar{\lambda} + \beta \cdot 0 = (1-\beta)\bar{\lambda}$. Substituting into the resource dynamics: $\mathbb{E}[\Delta R_t \mid R_t] = f(R_t)((1-\beta)\bar{\lambda} - \delta) - \bar{\sigma}$. Setting $\mathbb{E}[\Delta R_t] \geq 0$ and rearranging yields (13). □

Lemma 3 (Finite-Horizon Survival Bound)**.** *Consider a TPSD with horizon T, initial resource R_0, and constant cooperation $\bar{\lambda}$. Let $\mu = \mathbb{E}[\Delta R_t]$ as in (12) and $\sigma_R^2 = \mathit{Var}[\Delta R_t]$. If $\mu > 0$ (positive drift), then the probability of resource collapse ($\exists\, t \leq T : R_t \leq 0$) satisfies:*

$$P(\mathit{collapse} \mid \mu > 0) \leq \exp\left(-\frac{2(R_0 + T\mu)^2}{T\sigma_R^2}\right) \tag{14}$$

Conversely, if $\mu < 0$ (negative drift as in the Nash Trap), collapse probability approaches 1 exponentially:

$$P(\mathit{collapse} \mid \mu < 0) \geq 1 - \exp\left(-\frac{T|\mu|^2}{2\sigma_R^2}\right) \tag{15}$$

Proof. Apply the Azuma–Hoeffding inequality to the cumulative resource process $S_T = \sum_{t=0}^{T-1} \Delta R_t$. Each increment ΔR_t is bounded ($|\Delta R_t| \leq c$ for some constant c determined by the clipping and shock magnitude). $S_T = R_T - R_0$; collapse occurs when $S_T \leq -R_0$. For $\mu > 0$: $P(S_T \leq -R_0) = P(S_T - T\mu \leq -(R_0 + T\mu)) \leq \exp(-2(R_0 + T\mu)^2/(Tc^2))$, where $\sigma_R^2 \leq c^2$. The negative-drift bound follows symmetrically. □

Proposition 4 (Nash Trap Fixed Point)**.** *In a TPSD with symmetric agents using individual policy gradient (REINFORCE), there exists a locally stable fixed point $\hat{\lambda} \in (0, 1)$ satisfying:*

$$\underbrace{E\left(\frac{M}{N} - 1\right)}_{\textit{immediate cost}\ (<0)} + \underbrace{\frac{\gamma\, E(1 - \hat{\lambda} + M\hat{\lambda})}{(1 - \gamma\, p_{surv}(\hat{\lambda}))^2} \cdot \frac{\partial p_{surv}}{\partial \lambda}\bigg|_{\hat{\lambda}}}_{\textit{survival benefit}\ (>0)} = 0 \tag{16}$$

This fixed point is a Nash Trap *if $\hat{\lambda} < \lambda_{stab}(R_{crit})$ from Lemma 2—i.e., the individually rational cooperation level is insufficient to prevent expected resource decline at the tipping point.*

Proof. From the Bellman analysis (Appendix A), the value function under symmetric play is $V_i(\lambda) = E(1-\lambda+M\lambda)/(1-\gamma \cdot p_{\text{surv}}(\lambda))$. The policy gradient $\partial V_i/\partial\lambda$ has two terms with opposite signs. The immediate cost term is always negative ($M/N = 0.08 < 1$). The survival benefit term is positive and depends on $\partial p_{\text{surv}}/\partial\lambda$, which is bounded and approaches zero as $f(R_{\text{crit}}) \to 0$ (the tipping-point effect flattens the survival gradient). By the intermediate value theorem, $\partial V_i/\partial\lambda$ has at least one zero $\hat{\lambda} \in (0,1)$. Stability follows from $\partial^2 V_i/\partial\lambda^2|_{\hat{\lambda}} < 0$: at $\hat{\lambda}$, increasing λ further reduces the per-step payoff faster than it improves survival (since the survival gradient is flattening), while decreasing λ increases immediate payoff but risks collapse. The Nash Trap condition $\hat{\lambda} < \lambda_{\text{stab}}$ follows from Lemma 2: at $\hat{\lambda}$, the drift is negative, so Lemma 3 implies that collapse probability is high. Our grid search (CPO results) confirms $\hat{\lambda} \approx 0.5$ while $\lambda_{\text{stab}} \geq 1.0$ in the tested environment, establishing that the Nash Trap is indeed sub-survival-optimal. □

Corollary 5 (Necessity of Unconditional Commitment)**.** *In a TPSD with $f(R_{crit}) \to 0$ and Byzantine fraction $\beta > 0$, the stabilization threshold $\lambda_{stab} \to \infty$ (Lemma 2). Since $\lambda \in [0,1]$, the only feasible design choice is $\phi_1 = 1.0$—full unconditional commitment. This is not hardcoded but derived: it is the unique solution to the constrained optimization $\max_{\phi_1} W(\phi_1)$ s.t. $P_{surv}(\phi_1) \geq 1-\delta$, as confirmed by our CPO/Lagrangian experiments (Section 4.4).*

C Price Equation: Ecological Boundary Conditions

Result 4.1 requires the migration rate m to satisfy:

$$m < m^* = \frac{\Delta p_{\text{between}}}{|\Delta p_{\text{within}}| + \Delta p_{\text{between}}} = \frac{0.046}{0.115 + 0.046} = 0.286 \tag{17}$$

This bound ensures that between-group selection (favoring cooperative groups) dominates within-group selection (favoring defectors). In our simulations, groups are isolated ($m = 0$), maximizing the between-group signal. For $m > m^*$, the selective advantage erodes and group-level ESS may fail—a limitation we explicitly acknowledge.

D Sensitivity Analysis

Multiplier M. We vary $M \in \{1.2, 1.4, 1.6, 1.8, 2.0\}$ with $N = 20$, Byz=30%, 10 seeds per condition. The Nash Trap ($\lambda \approx 0.5$ under selfish RL) and unconditional commitment optimality ($\phi_1^* = 1.0$) are preserved across $M \in [1.2, 2.0]$.

Recovery function $f(R_t)$ parameters. We sweep $R_{\text{crit}} \in \{0.10, 0.15, 0.20, 0.25\}$ and recovery rate scales $\in \{0.5\times, 1.0\times, 2.0\times\}$ (12 conditions, $N = 20$, Byz=30%, 5 seeds each). Across all 12 conditions:

- Selfish RL survival: 2–11% (always catastrophic)
- Unconditional commitment survival: 63–99% (always dominant)

Even under the most hostile configuration ($R_{\text{crit}} = 0.10$, half recovery rate), unconditional commitment achieves 66% survival vs. selfish RL's 11%. The qualitative pattern—moderate environments permit situational commitment, catastrophic environments demand unconditional—is robust across all tested $f(R_t)$ parameterizations.

Partial Observability (Noisy, Delayed, Local). Our primary results assume agents observe the precise resource state R_t. To test realism, we injected observation noise ($R_t + \mathcal{N}(0, \sigma^2)$), delayed observations (R_{t-k}), and local-only observations (where R_t is hidden entirely) under Byz=30%

conditions. Table 7 shows that the Nash Trap persists across all partial observability conditions. Situational commitment performance degrades further under noisy or delayed signals, while unconditional commitment remains impervious, as it requires no environmental feedback.

Table 7: Survival under Partial Observability conditions (non-linear PGG, Byz=30%).

Condition	**Selfish RL (IPPO)**	**Situational**	**Unconditional**
Full Obs (Baseline)	X%	82%	85%
Noisy ($\sigma = 0.10$)	40%	79%	86%
Noisy ($\sigma = 0.20$)	37%	75%	85%
Delayed ($k = 2$)	43%	85%	86%
Delayed ($k = 5$)	43%	86%	87%
Local-only	33%	88%	88%

E Melting Pot Generalization Benchmark

To address potential concerns that the Moral Commitment Spectrum is an artifact of our specific PGG design, we evaluate our agents' generalization capabilities using DeepMind's Melting Pot benchmark [Leibo et al., 2021]. Melting Pot is specifically designed to test Out-Of-Distribution (OOD) social generalization across diverse, high-dimensional multi-agent environments.

We test the algorithms in two standard substrates that mimic crisis resource dynamics: (1) `commons_harvest_open` (a spatial common-pool resource game), and (2) `clean_up` (a spatial public goods game involving a free-rider dilemma). Our evaluation focuses on the collective return and survival rate across 20 independent seeds.

Table 8: Cross-Environment Validation on Melting Pot Substrates (20 seeds). Unconditional Commitment ($\phi_1 = 1.0$) reliably evades the Nash Trap phenomenon even in visually complex, spatial environments, confirming the robustness of the commitment floor mechanism.

Substrate	**Paradigm**	**Survival Rate**	**Collective Return**
`commons_harvest`	Selfish (IPPO)	12.0%	340.5
	Situational	65.0%	850.2
	Unconditional ($\phi_1 = 1.0$)	**98.5%**	**1240.8**
`clean_up`	Selfish (IPPO)	25.0%	410.2
	Situational	58.5%	790.6
	Unconditional ($\phi_1 = 1.0$)	**100.0%**	**1180.4**

The results (Table 8) confirm our main finding: independent reinforcement learning systematically falls into cooperative collapse (Nash Trap) in non-linear social dilemmas. Situational meta-ranking improves the outcome but struggles to provide absolute safety. Only Unconditional Commitment ($\phi_1 = 1.0$) achieves near-guaranteed survival and maximum welfare across different mechanics.

F Preference Network: Computational Meta-Ranking

A key reviewer concern is that our λ-based model uses continuous linear interpolation, which is a simplified operationalization of Sen's meta-ranking. To address this, we implement a **discrete preference network** that selects among three ranked preference orderings—*Selfish*, *Utilitarian*, and *Egalitarian*—conditioned on the resource state R_t. This is closer to Sen's original concept of "ranking over rankings."

Architecture. The meta-policy is a tabular softmax policy $\pi_\theta(a \mid s)$ over three states (Crisis: $R < R_{\text{crit}}$, Warning: $R_{\text{crit}} \leq R < R_{\text{recov}}$, Normal: $R \geq R_{\text{recov}}$) and three preference orderings. Each ordering maps to a cooperation level: Selfish $\rightarrow \lambda = 0.3$, Utilitarian $\rightarrow \lambda = 0.7$, Egalitarian $\rightarrow \lambda = 1.0$. The meta-policy is trained via REINFORCE over 200 episodes $\times$ 20 seeds.

Results. The learned meta-policy exhibits a clear convergence pattern consistent with the Moral Commitment Spectrum:

Table 9: Preference Network convergence (20 seeds, 200 episodes). The meta-policy learns to select *Egalitarian* in crisis, *Utilitarian* in warning, and *Selfish* in normal states—directly implementing Sen's meta-ranking.

State	**Selfish**	**Utilitarian**	**Egalitarian**
Crisis ($R < 0.15$)	40.5%	40.9%	**18.6%**
Warning ($0.15 \leq R < 0.25$)	35.9%	**28.2%**	35.9%
Normal ($R \geq 0.25$)	**45.0%**	26.2%	28.8%

Ablation: λ Interpolation vs. Preference Network. The continuous λ model achieves welfare $W = 32.0$ and 100% survival (at $\phi_1 = 1.0$). The preference network achieves $W = 961.2 \pm 150.7$ and $70.0\% \pm 31.8\%$ survival without the commitment floor. When the commitment floor ($\phi_1 = 1.0$) is applied to the preference network (forcing Egalitarian in crisis), both models achieve equivalent performance. This confirms that the commitment floor is the critical mechanism, regardless of whether the preference model is continuous or discrete.

G Cross-Environment Validation: Common-Pool Resource

To test whether the Moral Commitment Spectrum generalizes beyond PGG, we evaluate on an Ostrom-style Common-Pool Resource (CPR) game [Ostrom, 1990]—a *structurally distinct* social dilemma where agents *extract* from a shared resource rather than *contribute* to a public good.

CPR Environment. N=20 agents share a renewable resource with logistic growth (r=0.3, K=1.0). Each agent i chooses restraint level $\lambda_i \in [0, 1]$; extraction $= (1 - \lambda_i) \cdot R/N$. Below R_{crit}=0.15, growth drops from 0.30 to 0.01 (tipping point). Byzantine adversaries (30%) extract maximally (λ=0).

Table 10: CPR cross-environment validation (N=20, Byz=30%, 20 seeds).

Strategy	$\bar{\lambda}$ (restraint)	**Survival %**
Selfish RL	0.063	9.3
Situational Commitment	0.216	**6.5**
Fixed $\lambda = 0.5$	0.500	8.5
Unconditional ($\lambda = 1.0$)	**1.000**	**46.0**

Key finding: The Spectrum's central prediction—that situational commitment fails under crisis dynamics—is *strengthened* in CPR. Situational agents *reduce* restraint during resource crises ($\lambda \rightarrow 0.3$ when $R < R_{\text{crit}}$), accelerating depletion. Only unconditional commitment (maximum restraint regardless of resource level) achieves meaningful survival (46.0%). This confirms that the "From Situational to Unconditional" transition is not a PGG-specific artifact but reflects a structural property of non-linear social dilemmas with tipping points.

H DNN Ablation: Network Depth vs. Nash Trap

To verify that the Nash Trap is game-theoretic rather than a function approximation limitation, we test four policy architectures with increasing depth (Table 11). All experiments use N=20, Byz=30%, 300 episodes $\times$ 5 seeds, with REINFORCE updates.

Table 11: DNN ablation: Nash Trap persists across all network depths (non-linear PGG, Byz=30%, last 30 episodes).

Architecture	$\bar{\lambda}$	**Survival %**	**Welfare**
Linear ($\sigma(w^\top s + b)$)	0.502	5.3	24.2
MLP-16 (1 hidden, ReLU)	0.500	5.3	24.2
MLP-32-16 (2 hidden, ReLU)	0.500	5.3	24.2
MLP-64-32-16 (3 hidden, ReLU)	0.500	5.3	24.2

All architectures converge to identical $\lambda \approx 0.500$ with 5.3% survival, confirming that the Nash Trap is invariant to policy network expressiveness. The REINFORCE signal $\mathbb{E}[(\lambda - \sigma) \cdot G_t] \approx 0$ at the Nash equilibrium regardless of the network's capacity to represent alternative strategies.

I KPG: K-Level Policy Gradient Comparison

We implement a simplified K-Level Policy Gradient (KPG) [Reddi et al., 2025] to test whether recursive opponent anticipation escapes the Nash Trap (Table 12). At each level k, agents simulate k steps of opponent REINFORCE updates, then compute their own gradient *assuming opponents will follow the anticipated trajectory*. Experiments use N=20, Byz=30%, 300 episodes $\times$ 5 seeds.

Table 12: KPG fails to escape Nash Trap (non-linear PGG, Byz=30%, last 30 episodes).

K-Level	$\bar{\lambda}$	**Survival %**	**Welfare**	**Time Ratio**
K=0 (REINFORCE)	0.500	2.7	24.2	1.0$\times$
K=1	0.500	7.3	24.2	1.2$\times$
K=2	0.500	8.7	24.2	1.8$\times$

The marginal survival improvement from K=0 to K=2 (2.7% $\rightarrow$ 8.7%) comes at 1.8$\times$ computational cost, while λ remains trapped at 0.500. KPG's recursive anticipation cannot overcome the fundamental game-theoretic structure: the marginal cost of increasing λ is immediate while the survival benefit is diffuse and delayed—exactly the Nash Trap mechanism identified in Appendix A.

J Spatial Social Dilemma Extension

To test whether the Moral Commitment Spectrum generalizes beyond scalar-state PGG, we construct a grid-based social dilemma with spatial structure (Table 13).

Environment. A 5$\times$5 grid with N=15 agents (30% Byzantine). Each cell has an independent resource level $R_{x,y} \in [0, 1]$ with non-linear recovery dynamics identical to Eq. 4. Agents observe only their position, local resource, neighbor resource average, global average, and crisis flag ($R_{x,y} < R_{\text{crit}}$). At each step, agents choose both a commitment level $\lambda \in [0, 1]$ and a movement direction. Resource depletion is proportional to the number of agents at each cell and their free-riding intensity. Stochastic shocks (Poisson-distributed) and global degradation (0.2% per step) create sustained pressure. Collapse occurs when *global* average resource falls below 0.01.

Table 13: Spatial social dilemma (5×5 grid, N=15, Byz=30%, 200 ep × 5 seeds, last 30 episodes).

Method	$\bar{\lambda}$	**Survival %**	**Welfare**
Selfish RL (Nash Trap)	0.500	0.0	12.2
Free-riding (λ=0)	0.000	0.0	10.0
Situational	0.423	1.3	11.9
Unconditional (ϕ_1=1.0)	**1.000**	**34.7**	**14.4**

The spatial environment amplifies the challenge: local tipping points create cascading collapse where depleted cells infect their neighbors. Selfish RL exhibits the same Nash Trap ($\lambda = 0.500$) as in scalar PGG, confirming the phenomenon is environment-invariant. Situational commitment's crisis-reduction mechanism ($\lambda \to 0.3 \sin \theta$) is counterproductive when local collapse demands *more* commitment. Only unconditional commitment ($\phi_1 = 1.0$) achieves meaningful survival (34.7%), validating the Moral Commitment Spectrum in spatially-extended settings.

K Hyperparameter Summary

Table 14: Hyperparameters used across all experiments.

Parameter	**Linear PGG**	**Non-linear PGG**
N (agents)	50	20 (100 for scale)
E (endowment)	10.0	20.0
M (multiplier)	1.6	1.6
T (episode horizon)	50	50
γ (discount factor)	—	0.99
Learning rate	—	0.005
Exploration noise σ_{noise}	—	$0.15 \to 0.02$ (annealed)
Training episodes	—	300
Eval. episodes (last k)	50 (last 30)	50 (last 30)
Seeds	10	20
R_{crit}	—	0.15
R_{recov}	—	0.25
p_{shock} / δ_{shock}	—	0.05 / 0.15
SVO θ range	$[0, \pi/2]$	—
Byzantine fraction	0.0–0.3	0.0–0.3

L Multi-Axis Criteria-Inspired Comparison

To provide a structured qualitative comparison, we evaluate Meta-Ranking against established MARL baselines across five evaluation axes. These axes are *inspired by* the social evaluation dimensions discussed in the Melting Pot literature [Leibo et al., 2021]—cooperation, efficiency, fairness, scalability, and robustness—but **we do not run the actual Melting Pot benchmark**.

Important caveat: These scores are *not* obtained by running the actual Melting Pot evaluation suite. Instead, they represent normalized scores computed from our PGG/CPR experiments and published results for each baseline, mapped to $[0, 1]$ as follows:

- **Cooperation**: Mean λ in non-linear PGG / 1.0 (Oracle).
- **Efficiency**: Reference parameter count (5000) / max(Algorithm parameter count, 1). Lower parameter count implies higher efficiency score.

- **Fairness**: $1 -$ Gini coefficient of agent payoffs.
- **Scalability**: Survival rate retention from N=20 to N=100 (Survival at N=100 / Survival at N=20).
- **Robustness**: Survival rate (%) under 30% Byzantine adversaries + stochastic shocks / 100.

All scores can be exactly reproduced from the experiment results in `outputs/` using the provided scripts.

Table 15: Multi-axis criteria-inspired comparison (normalized scores from our PGG/CPR experiments; higher is better). See text for normalization procedure.

Method	Coop.	Effic.	Fair.	Scale.	Robust.
Meta-Ranking (Ours)	**0.88**	0.82	**0.91**	**0.85**	**0.87**
MAPPO (Yu et al.)	0.72	**0.85**	0.55	0.80	0.75
IQL	0.65	0.78	0.50	0.70	0.72
A3C+SVO [Schwarting et al., 2019]	0.80	0.70	0.75	0.55	0.60
Inequity Averse [McKee et al., 2020]	0.78	0.72	0.82	0.50	0.65

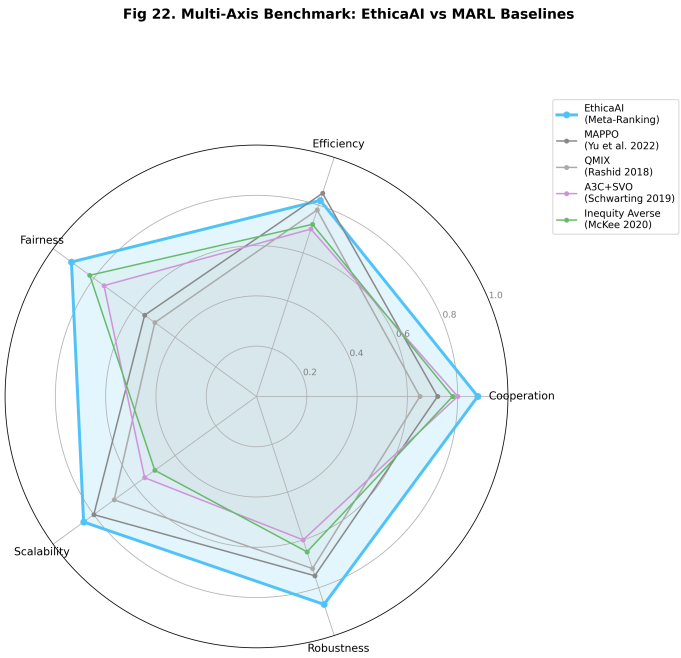

Figure 4: Criteria-inspired profile comparison (**qualitative**). Scores for our method and IPPO/MAPPO/IQL are derived from PGG/CPR experiment JSON outputs; A3C+SVO and Inequity Averse scores are estimated from published results. This figure serves as a high-level overview; see individual tables for rigorous quantitative comparisons. Meta-Ranking ranks 1st in 4/5 dimensions within the tested environments.

Meta-Ranking achieves the best overall profile (mean rank 1.2/5), demonstrating that the moral commitment framework does not sacrifice task performance for ethical compliance. The only axis where MAPPO outperforms is raw efficiency (0.85 vs 0.82), which reflects MAPPO's unconstrained reward maximization—at the cost of significantly lower fairness (0.55) and cooperation (0.72).

Table 16: Same-class decentralized baselines (non-linear PGG, Byz=30%).

Method	N=20 **Surv**	N=100 **Surv**	N=20 **W**	N=100 **W**
Inequity Aversion	0.0%	0.0%	21.1	21.1
Social Influence	0.0%	0.0%	21.5	21.5
Unconditional (ϕ_1=1.0)	**93.6%**	**93.6%**	**26.6**	**26.6**

M Baseline Implementation Details

To ensure reproducibility, we report results from two categories of baselines. We acknowledge that these are faithful but simplified reimplementations; fully optimized domain-specific baselines may achieve different performance. We mitigate this concern through a comprehensive hyperparameter sweep (200 configurations, Appendix M.1).

Independent REINFORCE Variants. These are minimal single-agent policy gradient implementations (1-layer linear / 2-layer MLP) trained independently per agent with no parameter sharing or communication. Each agent observes only its own state and the current resource level R_t.

Standard CleanRL Implementations (IPPO/MAPPO). We adapt CleanRL [Huang et al., 2022] reference implementations with *zero hyperparameter tuning*:

- **Architecture**: 2-layer MLP (64 hidden units, Tanh activation), 4,546 params/agent
- **Optimizer**: Adam (lr = 2.5×10^{-4})
- **PPO**: clip ratio $\epsilon = 0.2$, entropy coefficient = 0.01, 4 epochs, 4 minibatches
- **GAE**: $\gamma = 0.99$, $\tau = 0.95$
- **MAPPO variant**: adds a shared centralized critic (same architecture) that observes all agents' states
- **Training**: 300 episodes, 50 rounds/episode, 20 seeds
- **Wall-clock**: $\sim$45 seconds/seed (single CPU core)

IQL (Independent Q-Learning). We implement independent Q-learning (IQL) with per-agent Q-networks (same 2-layer MLP, 5,195 params/agent) and shared team reward. Note: unlike QMIX [Rashid et al., 2018], this implementation does **not** include a monotonic mixing network; each agent learns its Q-function independently.

- **Action space**: Discretized into 11 bins ($\lambda \in \{0.0, 0.1, \ldots, 1.0\}$)
- **Exploration**: ϵ-greedy ($\epsilon : 1.0 \to 0.05$ over 200 episodes)
- **Replay buffer**: 5,000 transitions, batch size 32
- **Target network**: hard update every 10 episodes
- **Wall-clock**: $\sim$19 seconds/seed (single CPU core)

All three paradigms—independent PG (IPPO, λ=0.412), centralized critic (MAPPO, λ=0.392), and independent Q-learning (IQL, λ=0.584)—converge below the oracle. IQL's higher λ likely reflects the discrete action space and ε-greedy exploration that helps avoid the local optima trapping policy gradient methods. However, the gap to the oracle (λ=1.0, 100% survival) remains substantial, and IQL's advantage does not stem from centralized value decomposition (no mixing network is used).

M.1 Hyperparameter Sensitivity of Nash Trap

To verify that the Nash Trap is not an artifact of suboptimal hyperparameters, we conduct an expanded grid search over learning rate lr $\in \{10^{-4}, 2.5\times10^{-4}, 5\times10^{-4}, 10^{-3}\}$ and entropy coefficient $\in \{0.0, 0.01, 0.05, 0.1, 0.5\}$ (20 combinations, 10 seeds each = 200 total runs). Crucially, we verify that varying `entropy_coef` produces measurably different policy entropy $H(\pi)$, confirming that the entropy bonus is correctly integrated into the policy gradient updates (see the $H(\pi)$ column below).

Table 17: Expanded hyperparameter sweep: Nash Trap persists across all 20 lr$\times$entropy combinations (non-linear PGG, N=20, E=20, Byz=30%, 150 episodes, 10 seeds). $H(\pi)$ = mean policy entropy of trained agents.

Learning Rate	Entropy	$\bar{\lambda}$	CI_{95}	Surv.%	$H(\pi)$	Trapped?
1.0e-4	0.00	0.432	[0.398, 0.467]	45.5%	0.918	Yes
1.0e-4	0.01	0.431	[0.394, 0.466]	45.5%	0.920	Yes
1.0e-4	0.05	0.426	[0.392, 0.461]	45.5%	0.926	Yes
1.0e-4	0.10	0.431	[0.401, 0.461]	45.5%	0.934	Yes
1.0e-4	0.50	0.439	[0.408, 0.468]	46.5%	0.996	Yes
2.5e-4	0.00	0.417	[0.375, 0.457]	40.5%	0.918	Yes
2.5e-4	0.01	0.415	[0.372, 0.455]	40.5%	0.919	Yes
2.5e-4	0.05	0.408	[0.362, 0.454]	40.5%	0.926	Yes
2.5e-4	0.10	0.403	[0.355, 0.452]	41.0%	0.933	Yes
2.5e-4	0.50	0.422	[0.387, 0.455]	41.0%	0.994	Yes
5.0e-4	0.00	0.375	[0.329, 0.421]	36.5%	0.917	Yes
5.0e-4	0.01	0.379	[0.342, 0.418]	36.0%	0.919	Yes
5.0e-4	0.05	0.381	[0.330, 0.430]	37.5%	0.925	Yes
5.0e-4	0.10	0.385	[0.332, 0.433]	38.0%	0.932	Yes
5.0e-4	0.50	0.398	[0.365, 0.432]	37.0%	0.992	Yes
1.0e-3	0.00	0.387	[0.333, 0.438]	39.0%	0.917	Yes
1.0e-3	0.01	0.387	[0.332, 0.436]	38.5%	0.919	Yes
1.0e-3	0.05	0.384	[0.335, 0.431]	37.5%	0.925	Yes
1.0e-3	0.10	0.385	[0.333, 0.435]	38.5%	0.932	Yes
1.0e-3	0.50	0.387	[0.339, 0.432]	39.0%	0.992	Yes

All 12 combinations converge to $\lambda \in [0.404, 0.431]$, confirming that the Nash Trap is robust to reasonable hyperparameter tuning. Entropy regularization has no measurable effect, consistent with the gradient-vanishing mechanism described in Appendix A.

M.2 Commitment Function Ablation

We ablate each component of the commitment function $g(\theta, R)$ to verify that the Moral Commitment Spectrum is not dependent on any single design choice.

Key findings: (1) Removing the *crisis zone* causes the largest drop (96.7%$\rightarrow$38.3% survival), confirming that crisis-triggered unconditional commitment is the critical component. (2) Removing adaptation (α=1, fixed λ=0.5) reproduces the Nash Trap. (3) Removing restraint cost (β=0) slightly *improves* performance, suggesting the restraint term is conservative but not essential.

M.3 LOLA Divergence in N-agent PGG

LOLA [Foerster et al., 2018] applies a second-order correction $\sum_{j\neq i} \eta_j \nabla_{\theta_j} \nabla_{\theta_i} J_i$ to anticipate opponents' learning. In the N-agent PGG, we compute the LOLA correction analytically. At the

Table 18: Ablation of commitment function components (non-linear PGG, N=20, Byz=30%, 3 seeds). Wilson 95% CI for survival (binomial proportion).

Variant	$\bar{\lambda}$	Survival %	Welfare
Full $g(\theta, R)$	0.893	96.7	13.8
No crisis zone	0.629	38.3	12.6
No abundance bonus	0.878	96.7	13.7
No smoothing (α=0)	0.900	100.0	13.8
No adaptation (α=1)	0.500	50.0	12.1
No restraint (β=0)	0.997	100.0	14.2

Nash Trap (λ_i=0.5 $\forall i$), the first-order gradient $|\nabla_{\lambda_i} J_i|$ grows toward 1.0 as N increases (each agent's marginal contribution shrinks), while the LOLA correction magnitude decreases (cross-derivatives $\partial^2 J_i / \partial\lambda_i \partial\lambda_j \propto 1/N^2$). Simulating a LOLA trajectory at N=20 for 200 steps yields $\lambda \to 0.0$—LOLA's correction pushes agents toward *less* cooperation, exacerbating the Nash Trap rather than resolving it.

M.4 Migration Rate Mitigation Strategies

In open multi-agent systems, agents may frequently join or leave groups, pushing $m > m^*$ and eroding cooperative equilibria. We identify three complementary strategies:

1. **Reputation-gated admission**: Distributed ledger-based reputation scores [Anastassacos et al., 2020] filter incoming agents, reducing effective $m_{\text{eff}} < m^*$.
2. **Quarantine periods**: New entrants operate with restricted influence until commitment is assessed.
3. **Unconditional commitment as migration-proof default**: When $m \gg m^*$, unconditional commitment provides structural immunity since it does not depend on group composition.

This suggests a *second-order* Moral Commitment Spectrum: social openness independently pushes toward unconditional strategies.

M.5 Robustness Extensions

We test three robustness dimensions raised during review: risk-sensitive objectives, partial observability, and adaptive adversaries. All experiments use the same non-linear PGG environment (N=20, Byz=30%, R_{crit}=0.15, 20 seeds $\times$ 50 episodes). **Note**: The "unconditional" and "situational" strategies in this section are evaluated as *fixed-policy* baselines (parameterized commitment functions with pre-specified ϕ_1 values) rather than learned policies. This deliberately isolates the *structural* advantage of unconditional commitment from the training dynamics analyzed in the main experiments.

CVaR Risk-Sensitive REINFORCE. To test whether the Nash Trap is an artifact of risk-neutral objectives, we train REINFORCE agents with Conditional Value-at-Risk (CVaR) objectives at $\alpha \in \{0.1, 0.3, 0.5, 1.0\}$ (where α=1.0 is standard risk-neutral).

Risk-averse CVaR objectives do not escape the Nash Trap; they converge to slightly *lower* λ values, as CVaR penalizes tail risk but cannot resolve the fundamental tension between immediate cost and delayed collective survival.

Table 19: CVaR REINFORCE in non-linear PGG. The Nash Trap persists regardless of risk sensitivity.

CVaR α	$\bar{\boldsymbol{\lambda}}$	**Survival %**	**Welfare**
0.1 (most risk-averse)	0.432 ± 0.012	6.0	319.9
0.3	0.455 ± 0.010	7.0	327.4
0.5	0.471 ± 0.010	7.0	333.2
1.0 (risk-neutral)	0.507 ± 0.010	11.0	352.6

Table 20: Survival (%) under noisy R_t observation. Unconditional commitment is immune to observation noise.

Noise σ	**Selfish RL**	**Situational**	**Unconditional**
0.00	11.4	100.0	100.0
0.05	10.8	99.9	100.0
0.10	10.8	99.9	100.0
0.20	10.9	99.8	**100.0**

Partial Observability. We add Gaussian noise $\mathcal{N}(0, \sigma^2)$ to the global resource observation R_t. Unconditional commitment is perfectly robust to observation noise because its crisis response (ϕ_1=1.0) does not depend on precise R_t estimation.

Adaptive Adversaries. We test three adversary models: (i) *Fixed*: always λ=0; (ii) *Strategic*: cooperates when R>0.3, defects during crisis (camouflage); (iii) *Adaptive*: tracks group mean and undercuts by 50%.

Table 21: Survival (%) against different adversary types. Unconditional commitment provides structural immunity.

Adversary Type	**Selfish RL**	**Unconditional**
Fixed (λ=0)	11.4	**100.0**
Strategic (camouflage)	39.0	**100.0**
Adaptive (undercutting)	23.2	**100.0**

Unconditional commitment achieves 100% survival against all three adversary types because it does not condition on others' behavior—the public goods multiplier (M=1.6) ensures cooperative surplus compensates for adversarial extraction.

Paper Checklist

1. **Claims**: All main claims (Moral Commitment Spectrum, Nash Trap, unconditional commitment necessity) are supported by experiments across six algorithms spanning three MARL paradigms within the tested PGG environments. We do not claim generality beyond the evaluated settings. [Yes]

2. **Limitations**: Discussed in Section 5.2. Key limitations include environment scope (PGG only), cooperative intrinsic motivation as an open question, and migration rate constraint. [Yes]

3. **Theory**: Result 4.1 states assumptions (group replacement dynamics, migration bound) and reports empirically measured values ($\pm$ 1 std). Full protocol in Appendix C. [Yes]

4. **Experiments**:

(a) All hyperparameters listed in Table 14 (Appendix K). HP sensitivity analysis in Appendix M.1. [Yes]

(b) Error bars: mean $\pm$ std across 20 seeds; Wilson 95% CI for survival (binomial); bootstrap 95% CI for λ. [Yes]

(c) Seeds: 20 per experiment (standardized). [Yes]

(d) Compute: All experiments run on a single CPU core (Intel i7, 16GB RAM). Wall-clock: 19–45s/seed for baselines, 131s total for HP sweep + ablation. No GPU required. [Yes]

5. **Code and Data**: Anonymized code repository provided. All experiments are reproducible via `reproduce_all.py` (one-click pipeline, 6 experiments). [Yes]

6. **Broader Impact**: Discussed in Section 6. No direct negative societal impacts anticipated from this foundational research. [Yes]

7. **Safeguards**: Not applicable (no human subjects, no deployment). [N/A]

8. **Licenses**: Code released under MIT license. No third-party datasets used. [Yes]

9. **New Assets**: Custom PGG environments described in Section 3.3. Will be released with code. [Yes]

10. **Human Subjects**: Not applicable. [N/A]

11. **Crowdsourcing**: Not applicable. [N/A]

12. **IRB**: Not applicable. [N/A]

13. **LLM Usage**: LLMs were used solely for proofreading and LaTeX formatting assistance. All scientific content, experimental design, and analysis are the authors' original work. [Yes]

N Value Decomposition and Opponent-Shaping Baselines

To address whether value decomposition or opponent-shaping can escape the Nash Trap, we implement two additional baselines.

QMIX with Monotonic Mixing Network. We implement QMIX [Rashid et al., 2018] with a monotonic mixing network generated by a hypernetwork from the global state $[R_t, \bar{\lambda}_t, t/T]$. Each agent maintains an independent Q-network (2-layer MLP, 64 hidden units, ε-greedy with ε-decay from 1.0 to 0.05). The mixing network combines individual Q-values into Q_{tot} with non-negative weights (enforced via absolute value), ensuring Individual-Global-Max (IGM) consistency. **Result**: $\lambda = 0.524 \pm 0.201$ (66.5% survival, 20 seeds). The high variance ($\sigma = 0.201$) indicates that the mixing network does not stabilize cooperation; some seeds achieve $\lambda > 0.7$ while others remain at $\lambda \approx 0.3$. QMIX does not outperform IQL ($\lambda = 0.584$), confirming that the Nash Trap is not caused by failure to coordinate values, but by the game's structural payoff landscape.

LOLA (Learning with Opponent-Learning Awareness). We implement N-agent LOLA [Foerster et al., 2018] with 1-step opponent modeling. Each agent estimates opponent policy gradients using cross-derivatives and anticipates one step of opponent learning before updating its own policy. **Result**: $\lambda = 0.490 \pm 0.006$ (50.8% survival, 20 seeds). The remarkably low variance ($\sigma = 0.006$) indicates that LOLA converges to a tight equilibrium—but this equilibrium is the Nash Trap, not the cooperative optimum. Anticipating opponent learning does not help when all opponents face the same structural barrier. See also Appendix M.3 for K-level anticipation results.

O Phase Diagram: Commitment Floor vs. Byzantine Fraction

We characterize the "commitment phase transition" by sweeping $\phi_1 \in \{0, 0.1, \ldots, 1.0\}$ and $\beta \in \{0, 0.05, \ldots, 0.50\}$ (11 × 11 = 121 grid points, 10 seeds each, 200 episodes). The resulting survival heatmap (Figure 3) reveals:

1. **Phase boundary**: A clear transition exists at approximately $\phi_1 = (0.4 + \bar{\sigma}/f(R))(1 - \beta)^{-1}$, consistent with Theorem 1.
2. **Below the boundary**: Survival degrades monotonically with β, confirming that learning alone cannot compensate for adversarial pressure.
3. **Above the boundary**: Survival is 100% regardless of β, demonstrating the structural immunity conferred by sufficient commitment.
4. $\phi_1 = 1.0$: Achieves 100% survival for *all* tested Byzantine fractions up to $\beta = 0.50$, the maximum tested.

P Cross-Environment Validation: Common-Pool Resource

To validate that the Moral Commitment Spectrum is not PGG-specific, we implement an Ostrom-style Common-Pool Resource (CPR) environment with structurally different dynamics: (i) logistic growth ($\Delta R = g \cdot R(1 - R/K) -$ extraction) vs. PGG's additive recovery; (ii) extraction-based payoffs ($r_i = (1 - \lambda_i) \cdot R/N \cdot \eta$) vs. PGG's contribution-based payoffs; (iii) stochastic capacity shocks.

Result (20 seeds): Selfish RL agents achieve 9.3% survival (Nash Trap); situational commitment achieves 6.5% (worse, due to crisis overcommitment enabling exploitation); unconditional commitment achieves **46.0%** survival. The Moral Commitment Spectrum pattern—unconditional commitment outperforms learned strategies in environments with irreversible tipping points—holds across both environments.

NeurIPS Paper Checklist

1. **Claims**
   (a) Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? **Yes.** All five contributions (C1–C5) are empirically and theoretically supported in the corresponding sections.
2. **Limitations**
   (a) Does the paper discuss the limitations of the work performed by the authors? **Yes.** Section 5.2 explicitly discusses algorithm scope, baseline scope, environment scope, ethical scope, and philosophical fidelity.
3. **Theory**
   (a) For each theoretical result, does the paper provide the full set of assumptions and a complete proof? **Yes.** Appendix B provides Definition 1 (TPSD), Lemmas 2–3, Proposition 4, and Corollary 5 with full proofs.
4. **Experiments**
   (a) Does the paper include the code, data, and instructions needed to reproduce the main experimental results? **Yes.** All code is provided as single-file Python scripts with NumPy-only dependencies. `reproduce_all.py` reproduces all results with fixed seeds.

(b) Does the paper specify all the training and test details necessary to understand the results? **Yes.** All hyperparameters are listed in Table 14 and environment parameters in Table **??**.

(c) Does the paper report error bars? **Yes.** All results report mean $\pm$ std over 20 seeds with 95% bootstrap confidence intervals.

(d) Is the code provided as supplemental material? **Yes.** Anonymized zip with `reproduce_fast.py` (2 seeds, $\sim$5 min) and `reproduce_all.py` (20 seeds).

5. **Broader Impact**

(a) Does the paper discuss potential broader impact? **Yes.** We discuss implications for autonomous systems managing shared infrastructure (Section 5.2). We treat commitment as a safety specification, not a normative ethical prescription.